

# 9. Beyond Traditional RDBMS

---

- ❖ Mega Trends in Data Management
- ❖ NoSQL (Not Only SQL) Databases
- ❖ NewSQL Databases

# 9. Beyond Traditional RDBMS

---

- ➡ Mega Trends in Data Management
- ❖ NoSQL (Not Only SQL) Databases
- ❖ New SQL Databases



# History of Data Management

- ❖ (Early-2000s) Business processing → Relational Database Management Systems (RDBMS)
  - ◆ E.g., Oracle, IBM DB2, Sybase, SQL Server, etc.
- ❖ (Mid-2000s) Internet blooming → low-cost RDBMS alternatives
  - ◆ Supported transactions, replication, recovery
  - ◆ Still must use custom middleware to scale out across multiple machines
  - ◆ Memcache for caching queries
  - ◆ E.g., MySQL, PostgreSQL

# History of Data Management



## ❖ (Late-2000s) Big Data → NoSQL

*“... the whole point of seeking alternatives is that you need to solve a problem that relational databases are a bad fit for ...”*

- Eric Evans

## ❖ Class of non-relational data storage systems

- ◆ Usually do not require a fixed table schema nor do they use the concept of joins

## ❖ Relax one or more of the ACID properties

- ◆ Brewer's CAP theorem

# Challenge 1: Scaling Up

---

- ❖ Datasets are just too big
- ❖ Hundreds of thousands of visitors in a short-time span → a massive increase in traffic
  - ◆ Developers begin to front RDBMS with a read-only cache to offload a considerable amount of the read traffic
  - ◆ Memcache or integrate other caching mechanisms within the application (i.e., Ehcache)
    - In-memory indexes, distributing and replicating objects over multiple nodes
  - ◆ As datasets grow, the simple memcache/MySQL model (for lower-cost startups) started to become problematic.



# Challenge 2: Availability


- ❖ A web-site is most likely to be unavailable when it is most needed → a huge volume of revenue loss
- ❖ Goal of web services today is to be as available as long as the network is on.
  - ◆ When some nodes crash or some communication links fail, the service still performs as expected
  - ◆ One desirable fault tolerance capability is to **survive a network partitioning into multiple parts.**
    - Distributed DBMSs (covered in the course) provides no solutions yet ...

## Traditional Approach

### *Structured & Repeatable Analysis*

#### **Business Users**

Determine what question to ask



#### **IT**

Structures the data to answer that question



Monthly sales reports  
Profitability analysis  
Customer surveys

## Big Data Approach

### *Iterative & Exploratory Analysis*

#### **IT**

Delivers a platform to enable creative discovery



#### **Business**

Explores what questions could be asked



Brand sentiment  
Product strategy  
Maximum asset utilization

# Possible Solutions to Scalability



- ❖ Began to look at multi-node database solutions
  - ◆ Distributed Database Systems
    - Basic principles and implementation techniques have been covered in the course
  - ◆ More techniques
    - To be covered by the next few slides



# Scaling RDBMS – Master/Slave

---

## ❖ Master-Slave

- ◆ All writes are written to the master. All reads performed against the replicated slave databases
- ◆ Critical reads may be incorrect as writes may not have been propagated down
- ◆ Large data sets can pose problems as master needs to duplicate data to slaves

## ❖ Multi-Master replication

# Scaling RDBMS - Partitioning

---

## ❖ Partition or sharding

- ◆ Scales well for both reads and writes
- ◆ Not transparent, application needs to be partition-aware (in contrast to DDB)
- ◆ Can no longer have relationships/joins across partitions
- ◆ Loss of referential integrity across shards

# Scaling RDBMS - NoSQL

---

- ❖ NoSQL systems are able to scale horizontally right out of the box:
  - ◆ Schemaless
  - ◆ Not ACID (i.e., eventual consistency)
  - ◆ Many are based on Google's Big Table or Amazon's Dynamo systems

# Scaling RDBMS - NewSQL

---

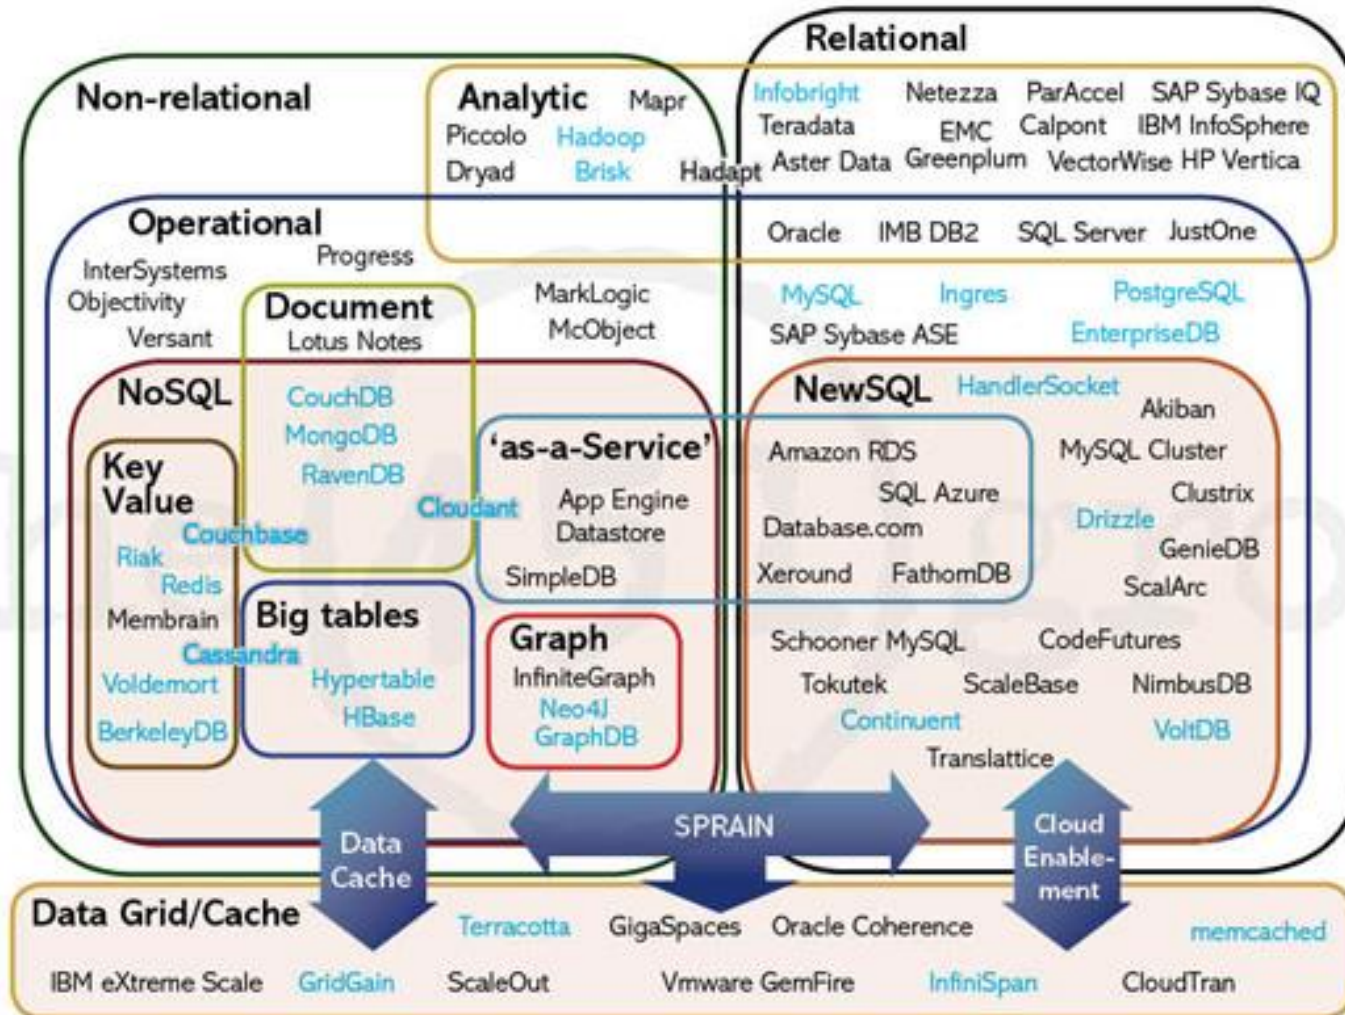
- ❖ (Early-2010s) Big Data → NewSQL
- ❖ New DBMSs that can scale across multiple machines natively and provide ACID guarantees
  - ◆ New Architectures
  - ◆ New MySQL storage engines
  - ◆ Transparent Clustering/Sharding

# NewSQL Definition

---

- ❖ SQL as the primary interface
- ❖ ACID support for transactions
- ❖ Non-locking concurrency control
- ❖ Higher per-node performance
- ❖ Parallel, shared-nothing architecture

# NoSQL, NewSQL, and Beyond (by 451 Group)



# 9. Beyond Traditional RDBMS

---

- ❖ Mega Trends in Data Management
- ☞ NoSQL (Not Only SQL) Databases
- ❖ New SQL Databases

# NoSQL (Not Only SQL)

---

- ❖ INSERT only, no UPDATE/DELETE
- ❖ No JOIN, thereby reducing query time
  - ◆ This involves de-normalizing data
- ❖ Lack of SQL support
- ❖ Non-adherence to ACID (Atomicity, Consistency, Isolation and Durability) properties



# Three Seeds of NoSQL

---

- ❖ BigTable (Google, 2006)
- ❖ Dynamo (Amazon, 2007)
  - ◆ Distributed key-value data store
- ❖ CAP Theorem (Eric A. Brewer)
  - ◆ BASE vs ACID

# The Perfect Storm

---

- ❖ Large datasets, acceptance of alternatives, and dynamically-typed data has come together in a perfect storm;
- ❖ Not a backlash/rebellion against RDBMS;
- ❖ SQL is a rich query language that cannot be rivaled by the current list of NoSQL (Not Only SQL) offerings.

# Google's BigTable



- ❖ A distributed storage system for managing structured data.
- ❖ Designed to scale to a very large size
  - ◆ Petabytes of data across thousands of servers
- ❖ Used for many Google projects
  - ◆ Web indexing, Personalized Search, Google Earth, Google Analytics, Google Finance, ...
- ❖ Flexible, high-performance solution for all of Google's products

# Motivation for BigTable

---

## ❖ Lots of (semi-)structured data at Google

- ◆ URLs:
  - Contents, crawl metadata, links, anchors, pagerank, ...
- ◆ Per-user data:
  - User preference settings, recent queries/search results, ...
- ◆ Geographic locations:
  - Physical entities (shops, restaurants, etc.), roads, satellite image data, user annotations, ...

## ❖ Scale is large

- ◆ Billions of URLs, many versions per page
- ◆ Hundreds of millions of users, thousands of queries per sec
- ◆ 100TB+ of satellite image data

# Why Not Just Use Commercial DB?

---

- ❖ Scale is too large for most commercial databases
- ❖ Even if it weren't, cost would be very high
- ❖ Low-level storage optimizations help performance significantly
  - ◆ Building internally means system can be applied across many projects for low incremental cost
  - ◆ Much harder to do when running on top of a database layer

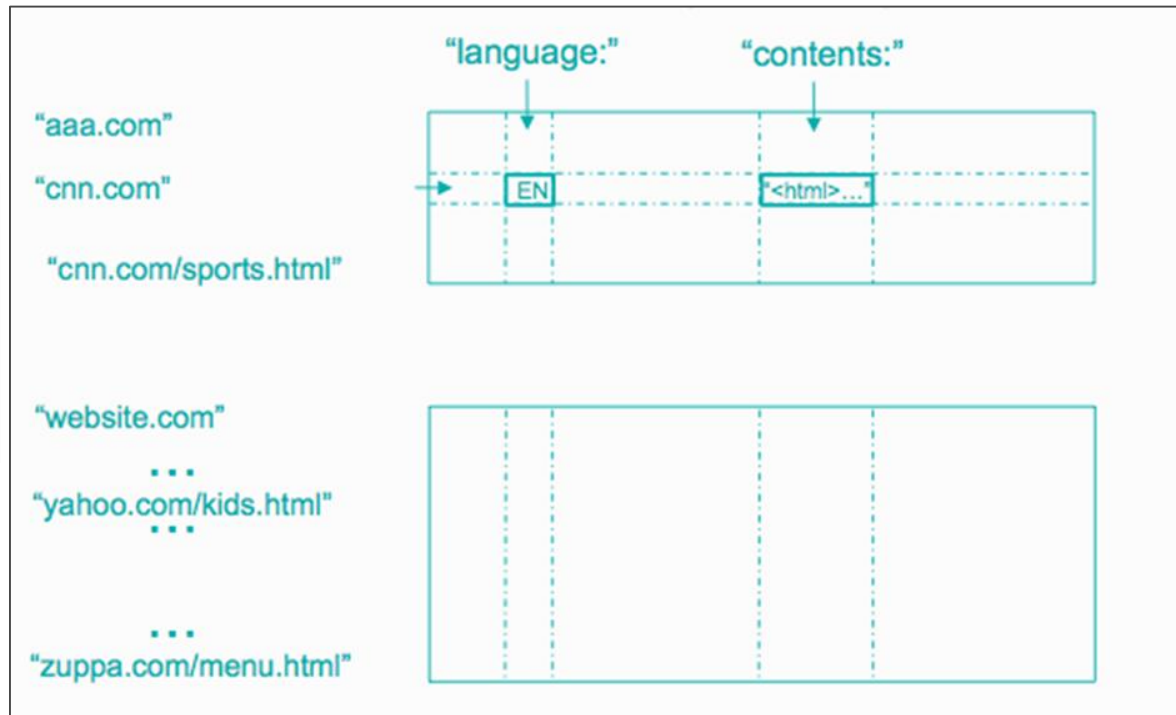

# Google's Goals

---

- ❖ Want asynchronous processes to be continuously updating different pieces of data
  - ◆ Want access to most current data at any time
- ❖ Need to support:
  - ◆ Very high read/write rates (millions of ops per second)
  - ◆ Efficient scans over all or interesting subsets of data
  - ◆ Efficient joins of large one-to-one and one-to-many datasets
- ❖ Often want to examine data changes over time
  - ◆ E.g. Contents of a web page over multiple crawls

# Basic Data Model - BigTable

- ❖ A sparse, distributed, persistent, multi-dimensional sorted map  
(row, column, timestamp) → cell contents



- ❖ Accommodate a large collection of web pages and related information

# Rows

---

- ❖ Use URLs as row keys
- ❖ Various aspects of a web page as column names
- ❖ Store contents of a web page in the **contents: column** under the timestamps when they were fetched.
- ❖ Name is an arbitrary string
  - ◆ Access to data in a row is atomic
  - ◆ Row creation is implicit upon storing data
- ❖ Rows ordered lexicographically
  - ◆ Rows close together lexicographically usually on one or a small number of machines



# Rows (*cont.*)

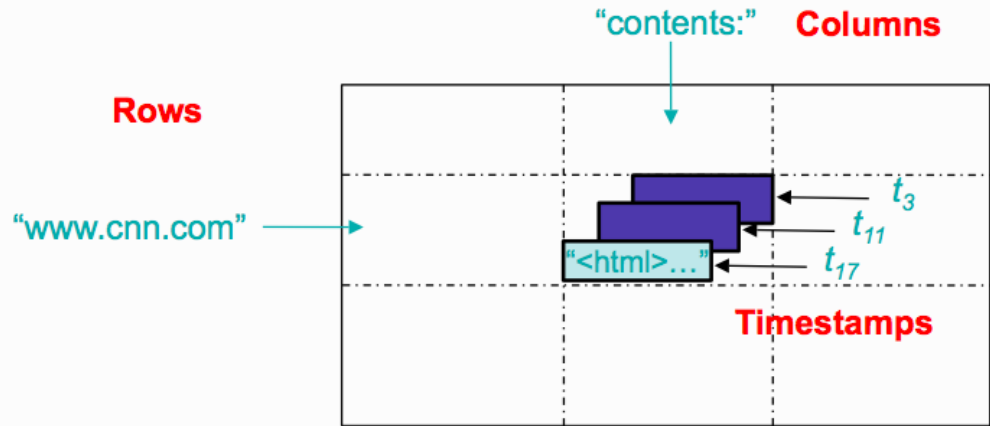
- ❖ Reads of short row ranges are efficient and typically require communication with a small number of machines.
- ❖ Can exploit this property by selecting row keys so they get good locality for data access.
- ❖ Example:

`math.gatech.edu, math.uga.edu, phys.gatech.edu,`  
`phys.uga.edu`

**VS**

`edu.gatech.math, edu.gatech.phys, edu.uga.math,`  
`edu.uga.phys`

# Timestamps



- ❖ Used to store different versions of data in a cell
  - ◆ New writes default to current time, but timestamps for writes can also be set explicitly by clients
- ❖ Items in a cell are stored in decreasing timestamp order.
- ❖ Application specifies how many versions of data items are maintained in a cell.
  - ◆ Bigtable garbage collects obsolete versions.

# BigTable API

---

## ❖ Implementation interfaces to

- ◆ create and delete tables and column families
- ◆ modify cluster, table, and column family metadata (e.g., access control rights)
- ◆ write or delete values in Bigtable
- ◆ look up values from individual rows
- ◆ iterate over a subset of the data in a table
- ◆ atomic R-M-W sequences on data stored in a single row key.

# Advantages of BigTable

---

- ❖ Distributed multi-level map
- ❖ Fault-tolerant, persistent
- ❖ Scalable
  - ◆ Thousands of servers
  - ◆ Terabytes of in-memory data
  - ◆ Petabyte of disk-based data
  - ◆ Millions of reads/writes per second, efficient scans
- ❖ Self-managing
  - ◆ Servers can be added/removed dynamically
  - ◆ Servers adjust to load imbalance

# BigTables in Google's Applications

Project name	Table size (TB)	Compression ratio	# Cells (billions)	# Column Families	# Locality Groups	% in memory	Latency-sensitive?
<i>Crawl</i>	800	11%	1000	16	8	0%	No
<i>Crawl</i>	50	33%	200	2	2	0%	No
<i>Google Analytics</i>	20	29%	10	1	1	0%	Yes
<i>Google Analytics</i>	200	14%	80	1	1	0%	Yes
<i>Google Base</i>	2	31%	10	29	3	15%	Yes
<i>Google Earth</i>	0.5	64%	8	7	2	33%	Yes
<i>Google Earth</i>	70	–	9	8	3	0%	No
<i>Orkut</i>	9	–	0.9	8	5	1%	Yes
<i>Personalized Search</i>	4	47%	6	93	11	5%	Yes

# Application 1: Google Analytics

---

- ❖ Enable webmasters to analyze traffic patterns at their web sites. Statistics such as:
  - ◆ Number of unique visitors per day and the page views per URL per day
  - ◆ Percentage of users that made a purchase given that they earlier viewed a specific page.
- ❖ How?
  - ◆ A small JavaScript program that the webmaster embeds in their web pages.
  - ◆ Every time the page is visited, the program is executed.
  - ◆ Program records information about each request:
    - user identifier and the pages being fetched

# Application 1: Google Analytics (*cont.*)

---

## ❖ Raw-Click BigTable (~ 200 TB)

- ◆ A row for each end-user session.
- ◆ Row name includes website's name and the time at which the session was created.
- ◆ Clustering of sessions that visit the same web site in a sorted chronological order.
- ◆ Compression factor: 6-7.

## ❖ Summary BigTable (~ 20 TB)

- ◆ Stores predefined summaries for each web site.
- ◆ Generated from the raw click table by periodically scheduled MapReduce jobs.
- ◆ Each MapReduce job extracts recent session data from the raw click table.
- ◆ Row name includes website's name and the column family is the aggregate summaries.
- ◆ Compression factor is 2-3.

# Application 2: Google Earth & Maps

---

- ❖ Move, view, and annotate satellite imagery at different resolution levels.
- ❖ One BigTable stores raw imagery (~ 70 TB):
  - ◆ Row name is a geographic segment. Names are chosen to ensure adjacent geographic segments are clustered together.
  - ◆ Column family maintains sources of data for each segment.
- ❖ There are different sets of tables for serving client data (e.g., index table).



# Application 3: Personalized Search

---

- ❖ Records user queries and clicks across Google properties.
- ❖ Users browse their search histories and request for personalized search results based on their historical usage patterns.

# Application 3: Personalized Search (*cont.*)

---

## ❖ One BigTable:


- ◆ Row name is userid
- ◆ A column family is reserved for each action type, e.g., web queries, clicks.
- ◆ User profiles are generated using MapReduce.
  - These profiles personalize live search results.
- ◆ Replicated geographically to reduce latency and increase availability.

- ❖ Huge Infrastructure
- ❖ Customer oriented business
- ❖ Reliability is key
- ❖ Guarantee Service Level Agreements
  - ◆ e.g., providing a response within 300ms for 99.9% of its requests for a peak client load of 500 requests per second.

# Amazon's Dynamo

## ❖ A distributed key-value storage system

- ◆ Simple
- ◆ Scale
- ◆ Highly available



	A	B	C	D	E	F	G	H
1		Attribute 1	Attribute 2	Attribute 3	...	Attribute <n>		
2	Item 1	value	value	value	value	value		
3	Item 2	value	value	value	value	value		
4	Item 3	value	value	value	value	value		
5	...	value	value	value	value	value		
6	Item <n>	value	value	value	value	value		
7								
8								
		Domain 1	Domain 2	...	Domain <n>			

# Requirements and Assumptions

---

## ❖ Query Model

- ◆ simple read and write operations to a data item that is uniquely identified by a key.

## ❖ ACID Properties

- ◆ Atomicity, Consistency, Isolation, Durability.

## ❖ Efficiency

- ◆ latency requirements which are in general measured at the 99.9th percentile of the distribution.

## ❖ Other Assumptions

- ◆ operation environment is assumed to be friendly and there are no security related requirements such as authentication and authorization.

# Amazon SimpleDB

---

- ❖ A web service based on Amazon Simple Storage Service (Amazon S3) and Amazon Elastic Compute Cloud (Amazon EC2)
- ❖ It stores, processes, and queries structured data in real time without operational complexity.
- ❖ It requires no schema, automatically indexes data, and provides a simple API for storage and access.
  - ◆ eliminating the administrative burden of data modeling, index maintenance, and performance tuning
- ❖ Developers gain access to its functionality within Amazon's computing environment, are able to scale instantly, and pay for what they use.

# Features of SimpleDB

---

- ❖ Simple to use
- ❖ Flexible
- ❖ Scalable
- ❖ Fast
- ❖ Reliable
- ❖ Inexpensive
- ❖ Designed for use with other Amazon Web services

# SimpleDB – Simple to Use

---

- ❖ Allowing users to quickly add data and easily retrieve or edit that data through a simple set of web service based API calls.
- ❖ Eliminating the complexity of maintaining and scaling users' operations.



# SimpleDB - Flexible

---

- ❖ Unnecessary to pre-define all of the data formats one will need to store; simply add new attributes to the data set when needed, and the system will automatically index the data accordingly.
- ❖ Storing structured data without first defining a schema provides developers with greater flexibility when building applications.

# SimpleDB - Scalable

---

- ❖ Allowing one to easily scale applications. Users can quickly create new domains as the data grows or your request throughput increases.
- ❖ Currently, users can store up to 10 GB per domain and can create up to 250 domains.

# SimpleDB - Fast

---

- ❖ Providing quick, efficient storage and retrieval of data to support high performance web applications.

# SimpleDB - Reliable

---

- ❖ The service runs within Amazon's high-availability data centers to provide strong and consistent performance.
- ❖ To prevent data from being lost or becoming unavailable, users' fully indexed data is stored redundantly across multiple servers and data centers.

# SimpleDB - Inexpensive

---

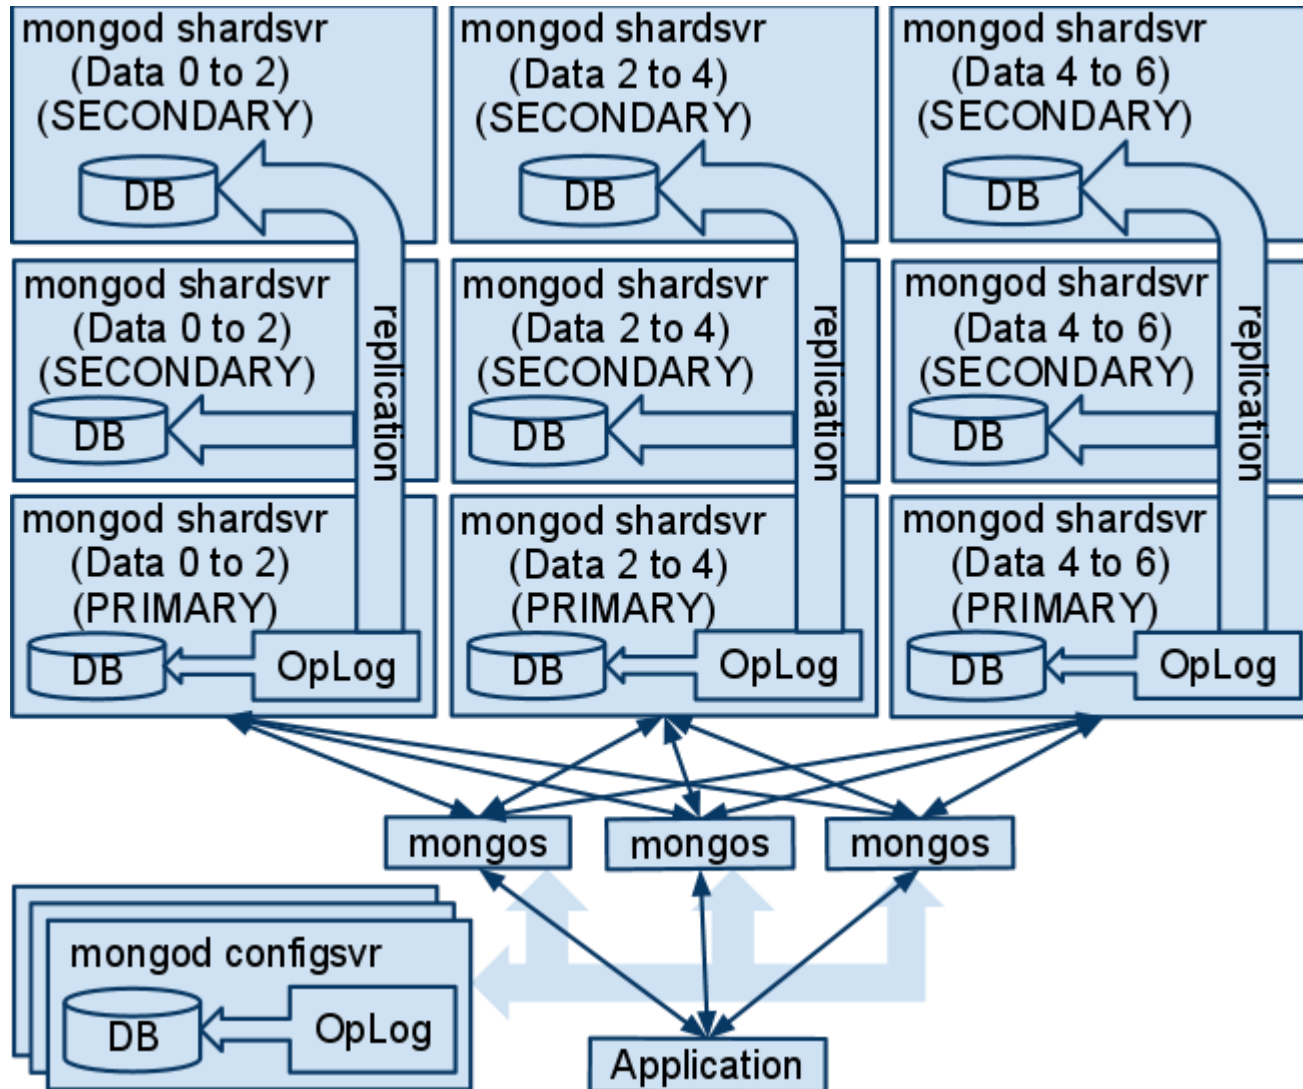
- ❖ Users pay only for resources they consume.
- ❖ Avoiding significant up-front expenditures traditionally required to obtain software licenses and purchase and maintain hardware, either in-house or hosted.
- ❖ Freeing users from many of the complexities of capacity planning, transforms large capital expenditures into much smaller operating costs, and eliminating the need to over-buy "safety net" capacity to handle periodic traffic spikes.

# SimpleDB – Integration with other Amazon Web Services

---

- ❖ Integrating with other Amazon web services such as Amazon EC2 compute cloud and Amazon S3 storage.
  - ◆ E.g., developers can query the object metadata from within the application in Amazon EC2 and return pointers to the objects stored in Amazon S3.

# MongoDB Architecture



- ❖ Easy to use
- ❖ Becoming more like a DBMS over time
- ❖ No transactions

# 9. Beyond Traditional RDBMS

---

- ❖ Mega Trends in Data Management
- ❖ NoSQL (Not Only SQL) Databases
- 👉 New SQL Databases



# NewSQL Solutions

---

- ❖ NewSQL is a class of database systems that aim to provide the same scalable performance of NoSQL systems while still maintaining the ACID guarantees of traditional SQL relational databases.
  - ◆ When the application needs to handle very large datasets or a very large number of transactions
  - ◆ When ACID guarantees are required
  - ◆ When the application can significantly benefit from the use of the relational model and SQL

# NewSQL Database Features

---

- ❖ SQL as the primary mechanism for application interaction.
- ❖ ACID support for transactions.
- ❖ A non-locking concurrency control mechanism so real-time reads will not conflict with writes.
- ❖ An architecture providing much higher per-node performance than available from traditional RDBMS solutions.
- ❖ A scale-out, shared-nothing architecture, capable of running on a large number of nodes without suffering bottlenecks.

# Categorization of NewSQL Solutions

---

- ❖ New Architectures
- ❖ New MySQL Storage Engines
- ❖ Transparent Clustering/Sharding

# New Architectures

---

- ❖ Newly designed from scratch to achieve scalability and performance (operate in a distributed cluster shared-nothing nodes)
- ❖ One of the key considerations in improving the performance is making non-disk (memory) or new kinds of disks (flash/SSD) the primary data store
- ❖ Some (hopefully minor) changes to the code
- ❖ Data migration is needed
- ❖ Solutions can be software-only (VoltDB, NuoDB and Drizzle) or supported as an appliance (Clustrix, Translattice).
- ❖ Examples: VoltDB, NuoDB, Clustrix, Drizzle, Translattice, and VMware's SQLFire

# New Database Example - VoltDB

- ❖ In-memory database
- ❖ ACID-compliant RDBMS
- ❖ Use a shared nothing architecture
- ❖ Written in Java and C++
- ❖ Supported operation systems:
  - ◆ Linux and Mac OS X
- ❖ Provides client libraries for
  - ◆ Java, C++, C#, PHP, Python and Node.js



# ACID in VoltDB

- ❖ **Atomicity:** VoltDB defines a transaction as a stored procedure, which either succeeds or rolls back on failure
- ❖ **Consistency:** VoltDB enforces schema and datatype constraints in all database queries
- ❖ **Isolation:** VoltDB transactions are globally ordered and run to completion on all affected partitions without interleaving
- ❖ **Durability:** VoltDB provides replication of partitions, and periodic database snapshots combined with command logging to ensure high availability and database durability

# New MySQL Storage Engines

---

- ❖ MySQL is used extensively in OLTP
- ❖ To overcome MySQL's scalability problems, a set of storage engines are developed
- ❖ Use the same programming interface as MySQL but scale better
- ❖ Examples: Xeround, Akiban, MySQL NDB cluster, GenieDB, Tokutek, etc.
- ❖ The good part is the usage of the MySQL interface, but the downside is data migration from other databases (including old MySQL) is not supported.

# Transparent Clustering/Sharding

---

- ❖ Provide a sharding middleware layer to automatically split databases across multiple nodes
- ❖ Retain the OLTP databases in their original format, but provide a pluggable feature to cluster transparently for scalability.
- ❖ Provide a transparent sharding middleware layer to automatically split databases across multiple nodes for scalability
- ❖ Both approaches allow reuse of existing skill sets and ecosystem, and avoid the need to rewrite code or perform any data migration.
- ❖ Examples: ScalArc, Schooner MySQL, dbShards and ScaleBase; and Continuent Tungsten




# H-Store

- ❖ An in-memory distributed database system
- ❖ Support for arbitrary transactions

## Two-Phase Commit

Transaction Prepare Request  
Transaction Prepare Response  
Transaction Finish Request  
Transaction Finish Response



# Google's NewSQL Solution


---

- ❖ Google's BigTable supports NoSQL and BASE
- ❖ Google also needs to support ACID
- ❖ Google's F1 is a hybrid database that combines high availability, the scalability of NoSQL systems like Bigtable, and the consistency and usability of traditional SQL databases.

# Google F1 and Spanner

❖ F1 is built on Spanner, which provides synchronous cross-data center replication and strong consistency.

- ◆ **Spanner: Google's Globally-Distributed Database**
- ◆ Synchronous replication implies higher commit latency
- ◆ Spanner mitigates this latency by using a **hierarchical schema model** with structured data types and through smart application design.




# F1's clustered hierarchical schema

	Traditional Relational	Clustered Hierarchical
Logical Schema	<p>Customer(<u>CustomerId</u>, ...)</p> <p>Campaign(<u>CampaignId</u>, CustomerId, ...)</p> <p>AdGroup(<u>AdGroupId</u>, CampaignId, ...)</p> <p>Foreign key references only the parent record.</p>	<p>Customer(<u>CustomerId</u>, ...)</p> <p>└─ Campaign(<u>CustomerId</u>, <u>CampaignId</u>, ...)</p> <p>    └─ AdGroup(<u>CustomerId</u>, <u>CampaignId</u>, <u>AdGroupId</u>, ...)</p> <p>Primary key includes foreign keys that reference all ancestor rows.</p>
Physical Layout	<p>Joining related data often requires reads spanning multiple machines.</p> <div><div>Customer(1,...) Customer(2,...)</div><div>AdGroup(6,3,...) AdGroup(7,3,...) AdGroup(8,4,...) AdGroup(9,5,...)</div></div> <div><div>Campaign(3,1,...) Campaign(4,1,...) Campaign(5,2,...)</div></div>	<div><div>Customer(1,...) Campaign(1,3,...) AdGroup (1,3,6,...) AdGroup (1,3,7,...) Campaign(1,4,...) AdGroup (1,4,8,...)</div><div>Related data is clustered for fast common-case join processing.</div></div> <p>Physical data partition boundaries occur between root rows.</p> <div>Customer(2,...) Campaign(2,5,...) AdGroup (2,5,9,...)</div>

# Besides

- ❖ **F1** also includes a fully functional distributed SQL query engine and automatic change tracking and publishing.



Google File System (GFS)

# F1 supports three types of transactions

---

- ❖ Each F1 transaction consists of multiple reads, optionally followed by a single write that commits the transaction.
- ❖ F1 implements three types of transactions, all built on top of Spanner's transaction support
  - ◆ Snapshot transaction
  - ◆ Pessimistic transaction
  - ◆ Optimistic transaction

# F1's Snapshot Transaction

---

- ❖ Read-only transaction with snapshot semantics
- ❖ Multiple client servers can see consistent views of the entire database at the same timestamp

# F1's Pessimistic Transaction

---

- ❖ The same as Spanner transactions
- ❖ Use a stateful communications protocol to require holding locks, so all requests in a single pessimistic transaction get directed to the same F1 server.
  - ◆ If the F1 server restarts, the pessimistic transaction aborts.
  - ◆ Reads in pessimistic transactions can request either shared or exclusive locks.



# F1's Optimistic Transaction

---

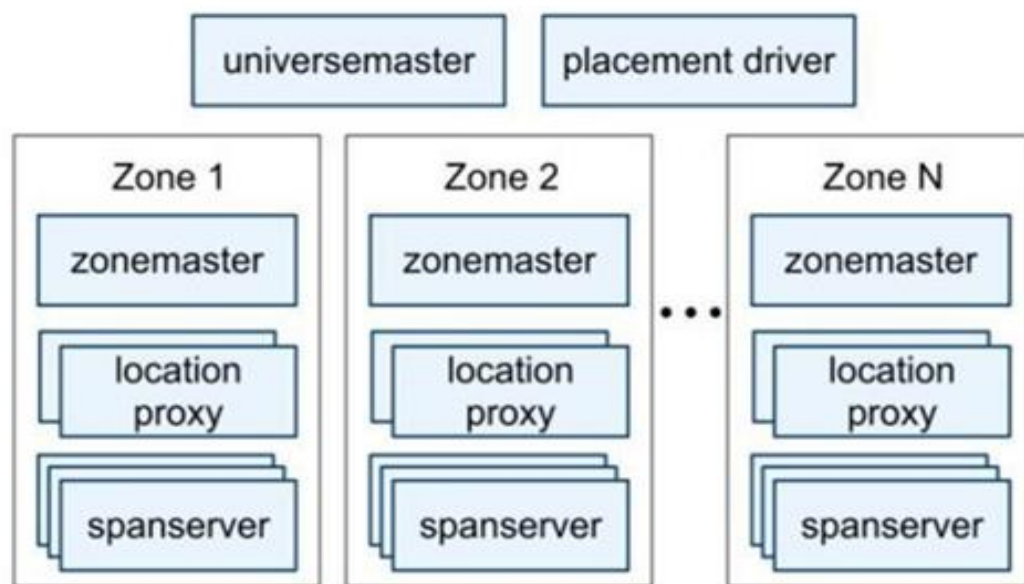
- ❖ Consist of an **arbitrarily long read** phase, which never takes Spanner locks, and then a **short write** phase.
- ❖ To detect row-level conflicts, F1 returns with each row its last modification timestamp, which is stored in a hidden lock column in that row.
- ❖ The new commit timestamp is automatically written into the lock column whenever the corresponding data is updated (in either pessimistic or optimistic transactions).

# F1's Optimistic Transaction (*cont.*)

---

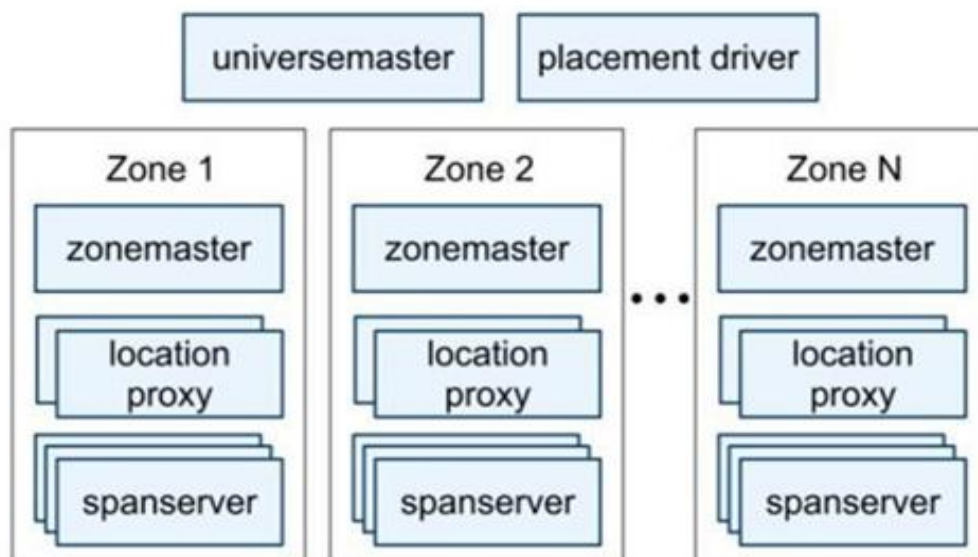
- ❖ The client library collects these timestamps, and passes them back to an F1 server with the write that commits the transaction.
- ❖ The F1 server creates a short-lived Spanner pessimistic transaction and re-reads the last modification timestamps for all read rows.
  - ◆ If any of the re-read timestamps differ from what was passed in by the client, there was a conflicting update, and F1 aborts the transaction;
  - ◆ Otherwise, F1 sends the writes on to Spanner to finish the commit.

# Architecture of Spanner



- ❖ A zone has one zonemaster and hundreds of spanservers.
  - ◆ zonemaster assigns data to spanservers
  - ◆ Spanservers serve data to clients.
  - ◆ The per-zone location proxies are used by clients to locate the spanservers for the requested data.

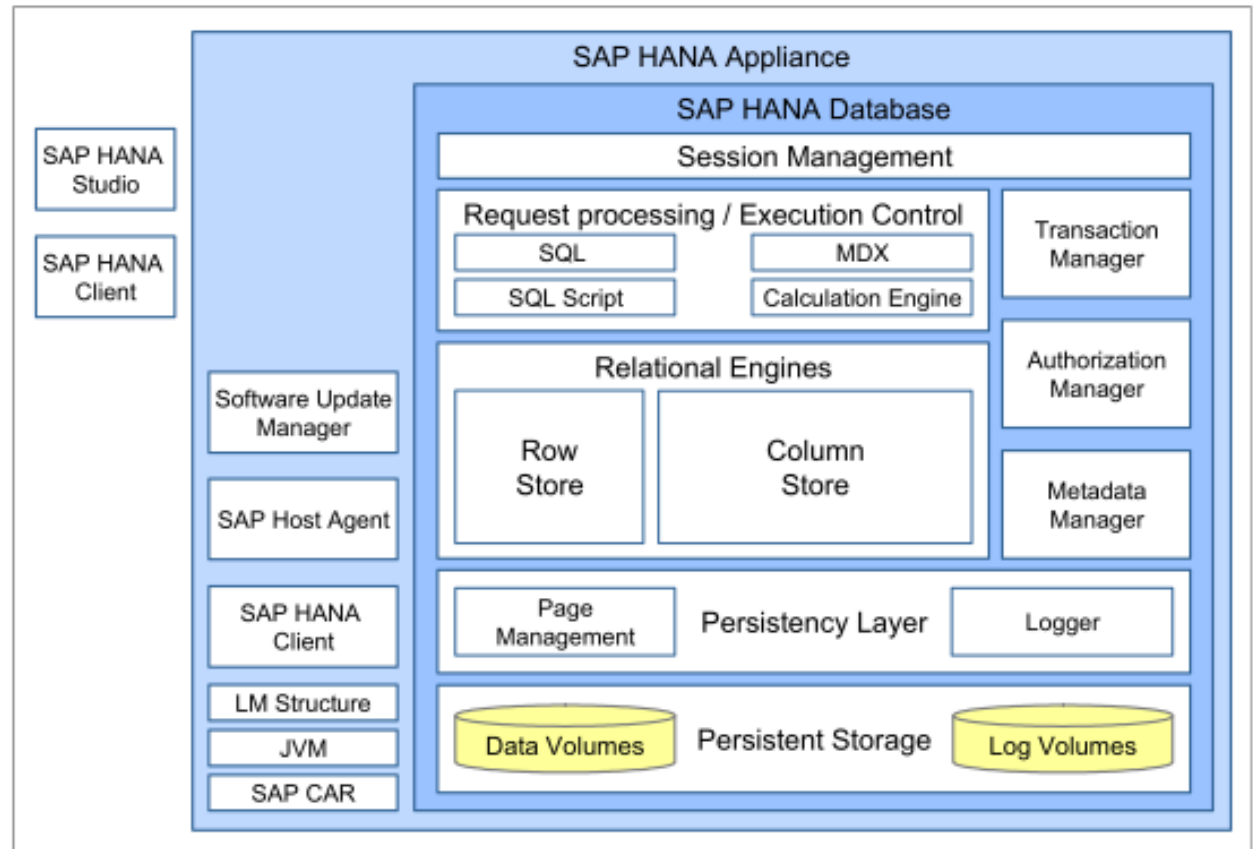
# Architecture of Spanner (cont.)



- ❖ The universe master displays status information about all the zones for interactive debugging.
- ❖ The placement driver handles automated movement of data across zones on the timescale of minutes.
  - ◆ It periodically communicates with the spanservers to find data that needs to be moved, either to meet updated replication constraints or to balance load.


# SAP Hana

- ❖ In-memory DB
- ❖ Combines row, column and object-oriented technology at the table level



# Summary

- ❖ The most powerful technologies take a while to mature. But when they do, they can rapidly retire mainstays that are decades old.



Gartner Inc.'s **hype cycle**: a graphic representation of the maturity, adoption, and social application of specific technologies

---

# Q & A