

# 5. Distributed Query Processing

---

Chapter 7

## Overview of Query Processing

Chapter 8

## Query Decomposition and Data Localization

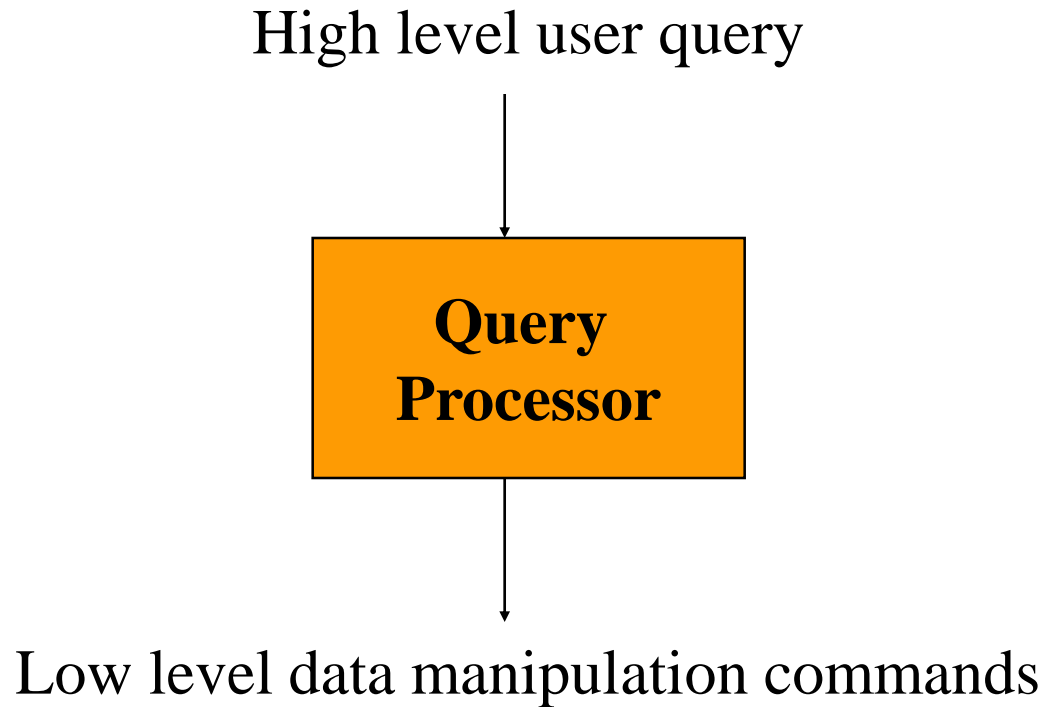
# Outline

---

- ❖ Overview of Query Processing (查询处理)
- ❖ Query Decomposition and Localization (查询分解与定位)

# Query Processing

---



# Query Processing Components

---

- ❖ Query language that is used
  - ◆ SQL (Structured Query Language)
- ❖ Query execution methodology
  - ◆ The steps that the system goes through in executing high-level (declarative) user queries
- ❖ Query optimization
  - ◆ How to determine the “best” execution plan?

# Query Language – Tuple Calculus

❖ Tuple calculus:  $\{ t \mid F(t) \}$

where  $t$  is a tuple variable, and  $F(t)$  is a well formed formula

❖ Example:

♦ *Get the numbers and names of all managers.*

$\{ t(ENO, ENAME) \mid t \in EMP \wedge t(TITLE) = "MANAGER" \}$

# Query Language – Domain Calculus

❖ Domain calculus:  $\{x_1, x_2, \dots, x_n \mid F(x_1, x_2, \dots, x_n)\}$

where  $x_i$  is a domain variable, and  $F(x_1, x_2, \dots, x_n)$  is a well formed formula

❖ Example:

$\{x, y \mid EMP(x, y, \text{"Manager"})\}$

Variables are position sensitive!

# Query Language SQL

---

❖ SQL is a tuple calculus language.

```
SELECT ENO, ENAME
```

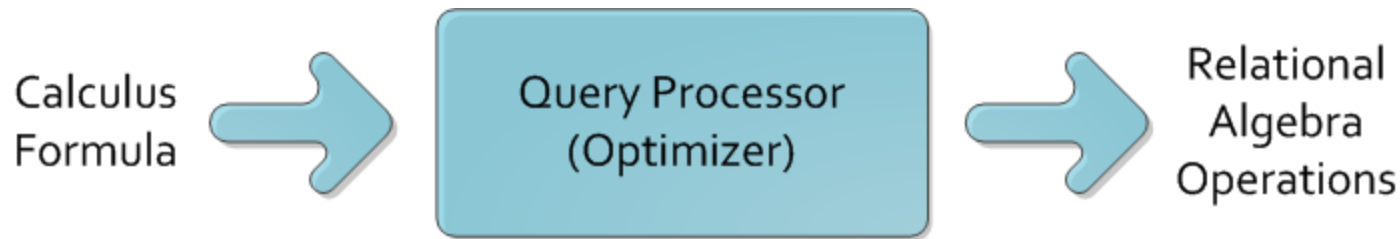
```
FROM EMP
```

```
WHERE TITLE="Programmer"
```

End user uses non-procedural (declarative) languages to express queries.

# Query Processing Objectives & Problems

- ❖ Query processor **transforms** queries into procedural operations to access data in an **optimal** way.



- ❖ Distributed query processor has to deal with query decomposition and data localization.



# Centralized Query Processing Alternatives

```
SELECT  ENAME
FROM    EMP E, ASG G
WHERE   E.ENO=G.ENO AND TITLE="manager"
```

❖ Strategy 1:  $\pi_{ENAME} \left( \sigma_{TITLE="manager" \wedge E.ENO=G.ENO} (E \times G) \right)$

❖ Strategy 2:  $\pi_{ENAME} \left( E \bowtie_{ENO} \sigma_{TITLE="manager"} (G) \right)$

Which one is better?

# Centralized Query Processing Alternatives (cont.)

```
SELECT  ENAME
FROM    EMP E, ASG G
WHERE   E.ENO = G.ENO AND TITLE="manager"
```

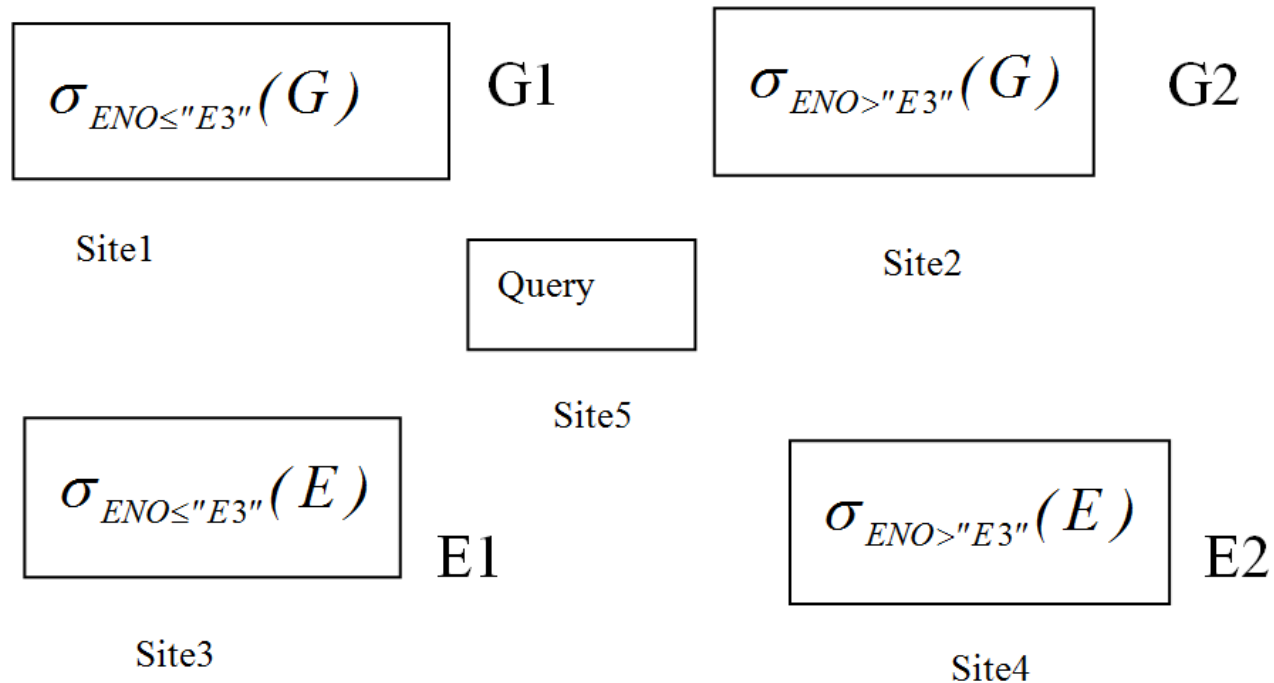
❖ Strategy 1:  $\pi_{ENAME} \left( \sigma_{TITLE="manager" \wedge E.ENO=G.ENO} (E \times G) \right)$

❖ Strategy 2:  $\pi_{ENAME} \left( E \bowtie_{ENO} \sigma_{TITLE="manager"} (G) \right)$

Strategy 2 avoids Cartesian product, so is “better”.

# Distributed Query Processing

- ❖ Query processor must consider the **communication cost** and **select the best site**.
- ❖ The same query example, but relation **G** and **E** are fragmented and distributed.



# Distributed Query Processing Plans

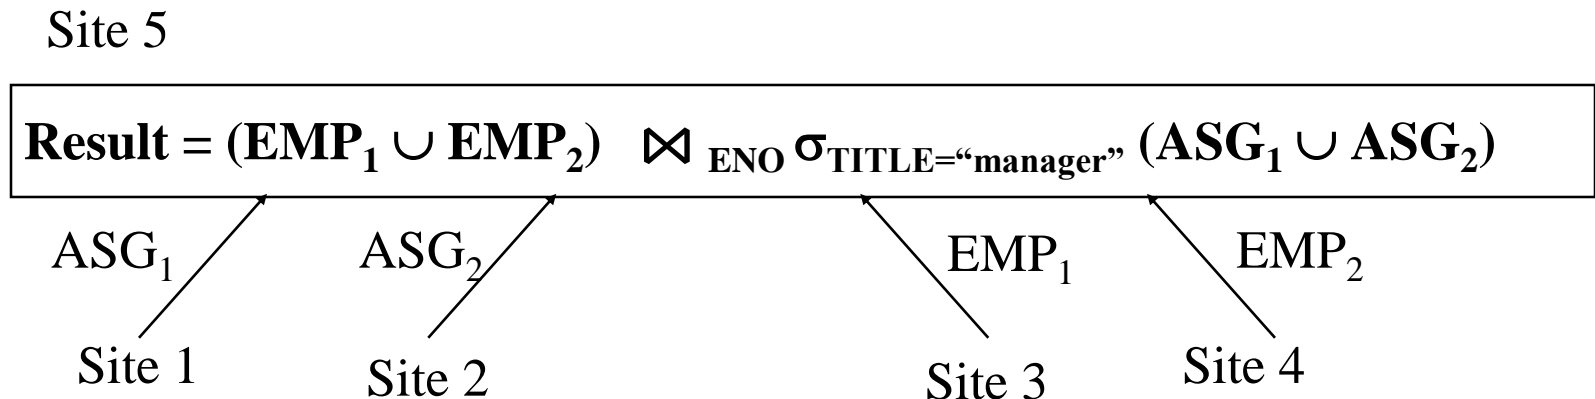
- ❖ By centralized optimization,

$$\pi_{ENAME} \left( E \bowtie_{ENO} \sigma_{TITLE="manager"}(G) \right)$$

- ❖ Two distributed query processing plans

# Distributed Query Plan I

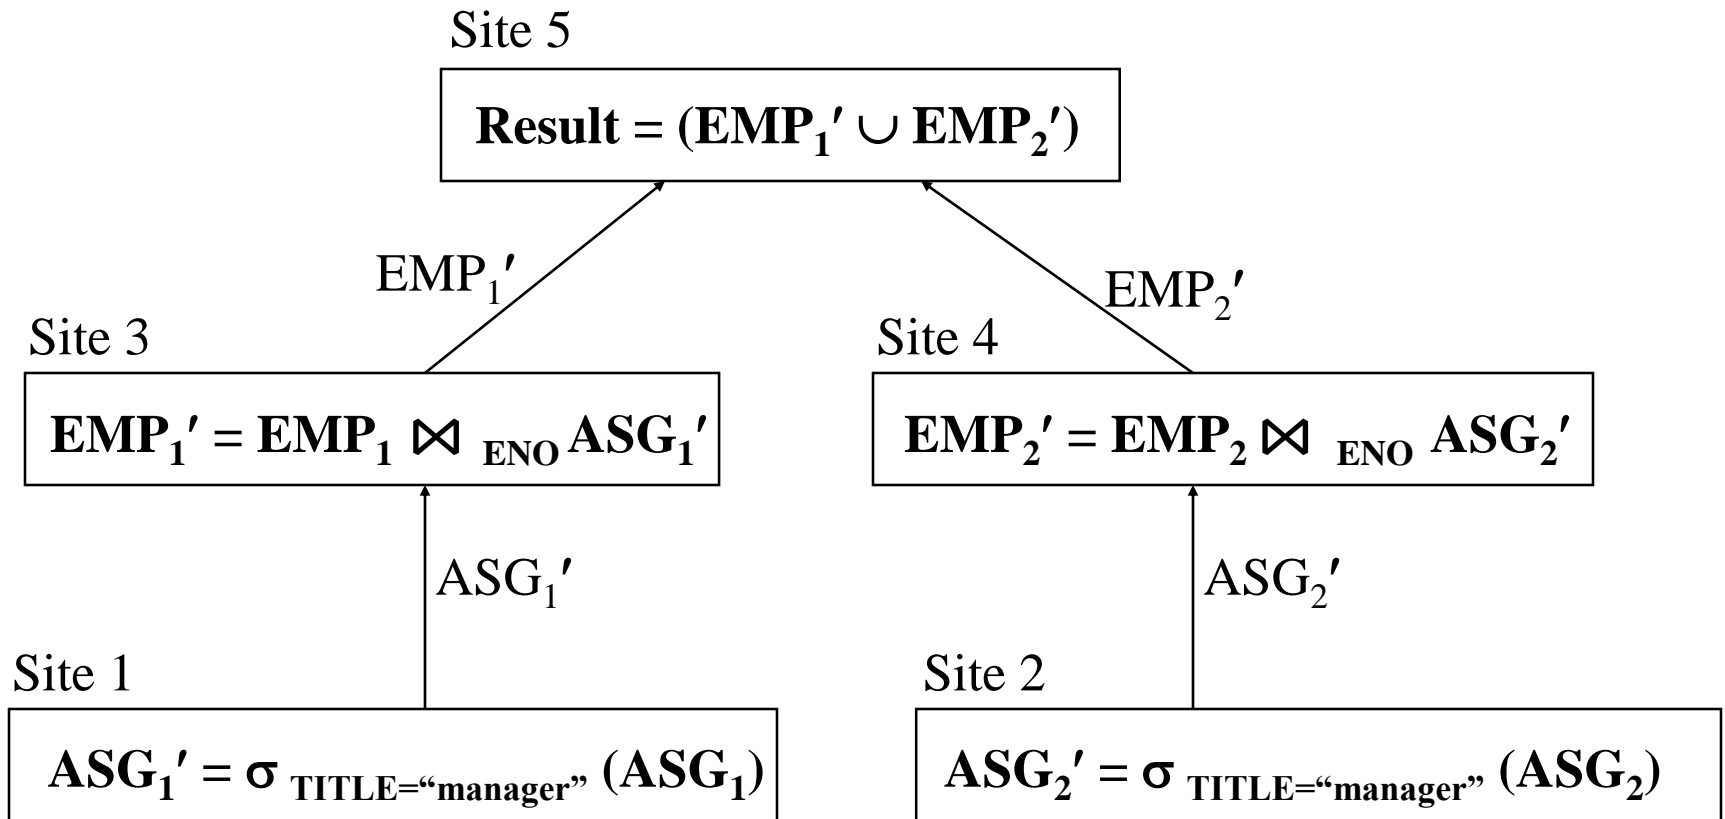
Plan I: To transport all segments to query site 5 and execute there.



This causes too much network traffic, very costly.

# Distributed Query Plan II

Plan II (Optimized):  $\pi_{ENAME} (E \bowtie_{ENO} \sigma_{TITLE="manager"} (G))$



# Costs of the Two Plans

$$\pi_{ENAME} \left( E \bowtie_{ENO} \sigma_{TITLE="manager"}(G) \right)$$

## ❖ Assume

- ♦  $size(EMP)=400$ ,  $size(ASG)=1000$ , 20 tuples with  $TITLE="manager"$
- ♦ tuple access cost = 1 unit; tuple transfer cost = 10 units
- ♦ ASG and EMP are locally clustered on attribute  $TITLE$  and  $ENO$ , respectively.

## ❖ Plan 1

- |  |             |
|--|-------------|
| ♦ Transfer EMP to site 5: $400 \times \text{tuple transfer cost}$      | 4000        |
| ♦ Transfer ASG to site 5: $1000 \times \text{tuple transfer cost}$     | 10000       |
| ♦ Produce ASG ' : $1000 \times \text{tuple access cost}$               | 1000        |
| ♦ Join EMP and ASG ' : $400 \times 20 \times \text{tuple access cost}$ | <u>8000</u> |
| Total cost   | 23,000      |

## ❖ Plan 2

- |   |            |
|---|------------|
| ♦ Produce ASG ' : $(10+10) \times \text{tuple access cost}$                       | 20         |
| ♦ Transfer ASG ' to the sites of EMP: $(10+10) \times \text{tuple transfer cost}$ | 200        |
| ♦ Produce EMP ' : $(10+10) \times \text{tuple access cost}$                       | 20         |
| ♦ Transfer EMP ' to result site: $(10+10) \times \text{tuple transfer cost}$      | <u>200</u> |
| Total cost  | 440        |

# Query Optimization Objectives

---

- ❖ Minimize a cost function

I/O cost + CPU cost + communication cost

- ♦ These might have different weights in different distributed environments
- ♦ Can also maximize throughput



# Communication Cost

---

## ❖ Wide area network

- ◆ Communication cost will dominate
  - Low bandwidth
  - Low speed
  - High protocol overhead
- ◆ Most algorithms ignore all other cost components

## ❖ Local area network

- ◆ Communication cost not that dominate
- ◆ Total cost function should be considered

# Complexity of Relational Algebra Operations

- ❖ Measured by cardinality  $n$  and tuples are sorted on comparison attributes

Operation	Complexity
$\sigma$ , $\pi$ (without duplicate elimination)	$O(n)$
$\pi$ (with duplicate elimination), GROUP	$O(n \log n)$
<i>Join, Semijoin, Division</i> , <span style="background-color: #FFDAB9;"> </span> –	$O(n \log n)$
<i>Cartesian-Product</i> X	$O(n^2)$

# What is O-Notation?

---

- ❖ Algorithm Analysis Basics

- ❖ O-Notation:

Intuition, Definition, Manipulation, and Limitations

# Algorithm Analysis Basics

---

- ❖ Goal: measure the efficiency of algorithms.
  - ◆ How much time and space does an algorithm use when it processes an input?
- ❖ Question: how to characterize and measure the performance of an algorithm?
- ❖ Why not measure execution time of an algorithm?
  - ◆ Varying significantly with different computers, programming languages, and compilers

# Resource Consumption Patterns

❖ Observation: algorithms usually take more time/ space as the size of the problem grows

◆ A problem size  $n$  can be

- the length of a list that an algorithm searches;
- the number of nodes in a tree that an algorithm prints;
- the number of items in an array that an algorithm sorts, etc.

Array Size $n$	Computer1	Computer2
125	12.5	2.8
250	49.3	11.0
500	195.8	43.4
1000	780.3	172.9
2000	3114.9	690.5

Table 1: SelectionSort algorithm running times example in milliseconds on two types of computers.

# Measurement Example

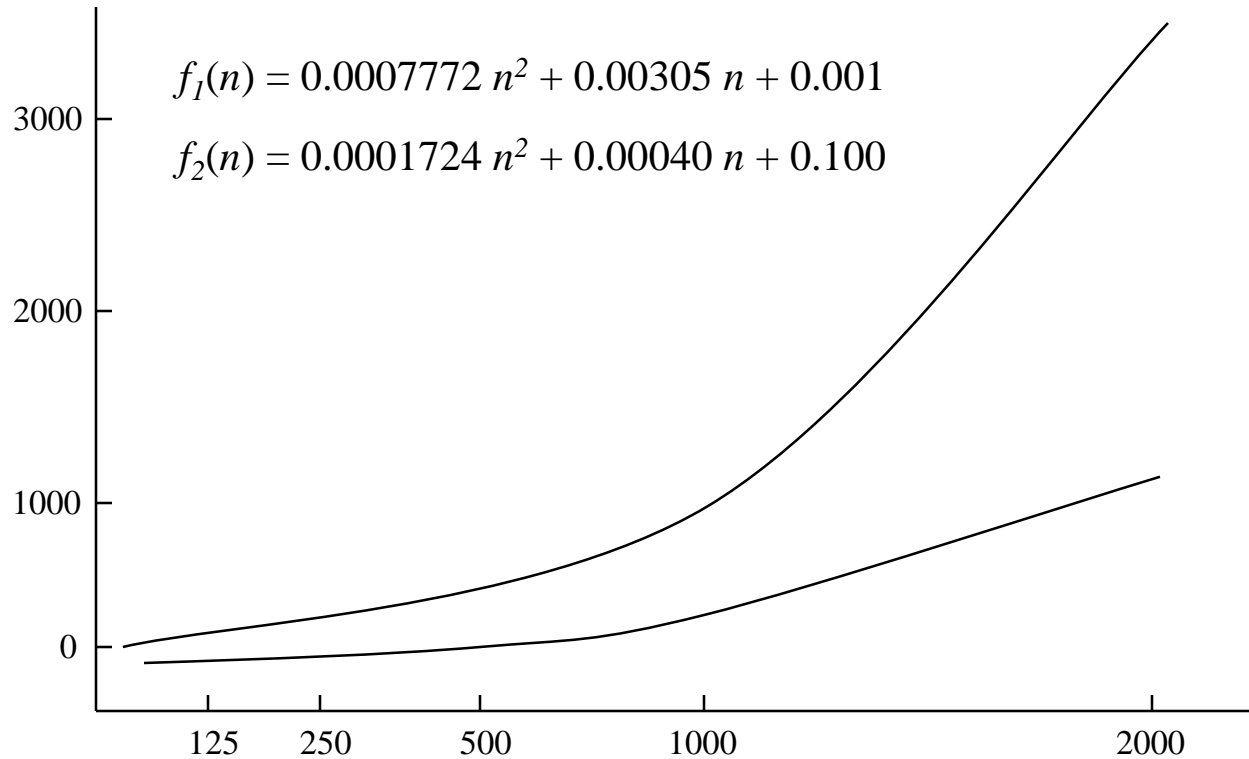
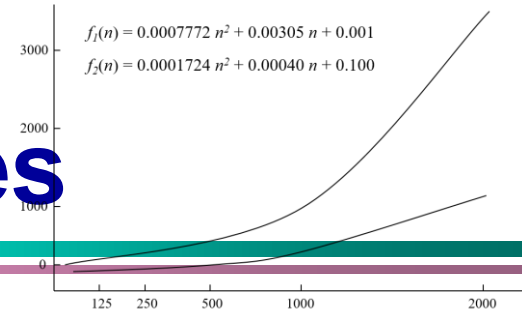


Figure 1: Two curves fitting the data in Table 1.

# Measurement Example Notes

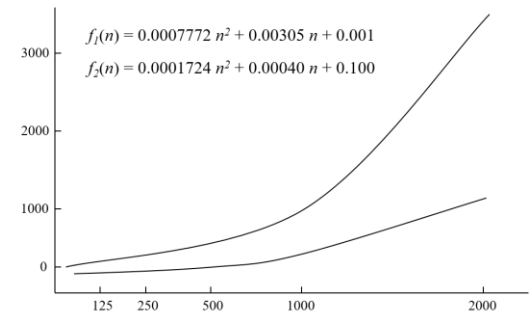


- ❖ The quadratic functions (the highest power is 2) shown are the only good fit for the data points gathered.
- ❖ The curves differ in their coefficients, but both are quadratic in terms of the size  $n$ .
- ❖ If continuing gathering execution times using different computers, languages, and compilers, we'll discover that the amount of time the algorithm consumes is best approximated by a quadratic curve of the form  $f(n) = an^2 + bn + c$ .

The *shape* of the curve depends only on the algorithm used.

# Measurement Example Generalization

- ❖ Each curve expresses an algorithm's resource consumption in terms of the problem size.
- ❖ Each algorithm is associated with a family of similarly shaped curves.
- ❖ Running times for different algorithms fall into different complexity classes, characterized by different families of curves.



Ignore coefficients and concentrate on the *kind* of function.



# Simplification of Complexity Functions

- ❖ Ignore the less terms and focus on the dominant one.
  - ◆ The dominant term is the one that grows faster when  $n$  grows

$f(n) = an^2 + bn + c$ where $a=0.0001724$ , $b=0.0004$ , $c=0.1$			
$n$	$f(n)$	$an^2$	$n$ -term as % of total
125	2.8	2.7	94.7
250	11.0	10.8	98.2
500	43.4	43.1	99.3
1000	172.9	172.4	99.7
2000	690.5	689.6	99.9

Table 2: Percentage contribution of  $n^2$  term to the total.

- ❖ Ignore the coefficient of the dominant term

# Intuition of O-Notation

## ❖ Two simplification rules lead to the O-notation

- ◆ e.g., the O-notation for the quadratic function,  $f(n)=an^2+bn+c$ , is  $O(n^2)$ .

$f(n)$	$O(f(n))$
$0.3n^2 + 20n + 512$	
$0.0001n^4 + 10000n^2$	
$42\log_2 n$	
$7n\log_{10} n + 2n - 12$	


O-notation gives us a language to describe algorithmic complexity

# Common Complexity Classes

Adjective Name	O-Notation
Constant	$O(1)$
Logarithmic	$O(\log n)$
Linear	$O(n)$
$n \log n$	$O(n \log n)$
Quadratic	$O(n^2)$
Cubic	$O(n^3)$
Exponential	$O(2^n)$
Exponential	$O(10^n)$

# Running Times Example

	2 ( $2^1$ )	16 ( $2^4$ )	256 ( $2^8$ )	1024 ( $2^{10}$ )	1048576 ( $2^{20}$ )
1	1 $\mu$ s	1 $\mu$ s	1 $\mu$ s	1 $\mu$ s	1 $\mu$ s
$\log_2 n$	1 $\mu$ s	4 $\mu$ s	8 $\mu$ s	10 $\mu$ s	20 $\mu$ s
n	2 $\mu$ s	16 $\mu$ s	256 $\mu$ s	1.02 ms	1.05 s
$n \log_2 n$	2 $\mu$ s	64 $\mu$ s	2.05 ms	10.2 ms	21 s
$n^2$	4 $\mu$ s	25.6 $\mu$ s	65.5 ms	1.05 s	1.8 wks
$n^3$	8 $\mu$ s	4.1 ms	16.8 s	17.9 min	36559 yrs
$2^n$	4 $\mu$ s	65.5 msec	$3.7 \cdot 10^{63}$ yrs	$5.7 \cdot 10^{294}$ yrs	$2.1 \cdot 10^{315639}$ yrs


 Estimated  
lifetime of  
the sun:  
 $5 \cdot 10^9$  yrs !

1 $\mu$ s = $10^{-6}$ s	1 s = one second	1 wk = 604800 s
1 ms = $10^{-3}$ s	1 min = 60 s	1 yr = 31557600 s

# Observations from Running Time Example

---

- ❖ For problems of small size (e.g.,  $n \leq 16$ ), the complexity class of  $f(n)$  does not matter much.
- ❖ For problems of medium size (e.g.,  $n=1024$ ), algorithms that are no more complex than  $n^2$  are still useful.
- ❖ For problems of large size (e.g.,  $n=1048576$ ), the difference between  $n \log_2 n$  and  $n^2$  algorithms is huge.
- ❖ Exponential algorithms (e.g.  $2^n$ ) tends to take a disastrously long time for all but small problems.

# Change of Observation Angle

Table 5: Size of the largest problem that an algorithm can solve if solution is computed in time  $\leq T$  at 1 microseconds per step.

$f(n) \backslash T$	1 min	1 hr	1 day	1 wk	1 yr
n	$6 \cdot 10^7$	$3.6 \cdot 10^9$	$8.64 \cdot 10^{10}$	$6.05 \cdot 10^{11}$	$3.15 \cdot 10^{13}$
$n \log_2 n$	$2.8 \cdot 10^6$	$1.3 \cdot 10^8$	$2.75 \cdot 10^9$	$1.77 \cdot 10^{10}$	$7.97 \cdot 10^{11}$
$n^2$	$7.75 \cdot 10^3$	$6.0 \cdot 10^4$	$2.94 \cdot 10^5$	$7.78 \cdot 10^5$	$5.62 \cdot 10^6$
$n^3$	$3.91 \cdot 10^2$	$1.53 \cdot 10^3$	$4.42 \cdot 10^3$	$8.46 \cdot 10^3$	$3.16 \cdot 10^4$
$2^n$	25	31	36	39	44
$10^n$	7	9	10	11	13

# Algorithms of Different Complexity

---

## ❖ Constant time algorithms $O(1)$

- ◆ Take no more than a *fixed* amount of time to run regardless of the problem size.
- ◆ e.g., choose and print a single random array item  $A[i]$  in an array  $A[0:n-1]$ .

## ❖ Linear time algorithms $O(n)$

- ◆ Run in time proportional to the problem size.
- ◆ e.g., search for occurrences of a give word in a document.

# Algorithms of Different Complexity (*cont.*)

---

## ❖ Exponential time algorithms $O(2^n)$

- ◆ Not practical to use for any but small problems.
- ◆ e.g., traveling salesperson problem solution, computing moves in game-playing situations.

## ❖ Quadratic and cubic time algorithms $O(n^2)$ / $O(n^3)$

- ◆ e.g., computation on an  $n \times n$  matrix, or in 3-dimensional space.



# Efficiency of List Implementation

List Operation	Sequential Representation	Linked representation
Find the length	$O(1)$	$O(n)$
Insert a new first item	$O(n)$	$O(1)$
Delete the last item	$O(1)$	$O(n)$
Replace the $i^{\text{th}}$ item	$O(1)$	$O(n)$
Delete the $i^{\text{th}}$ item	$O(n)$	$O(n)$

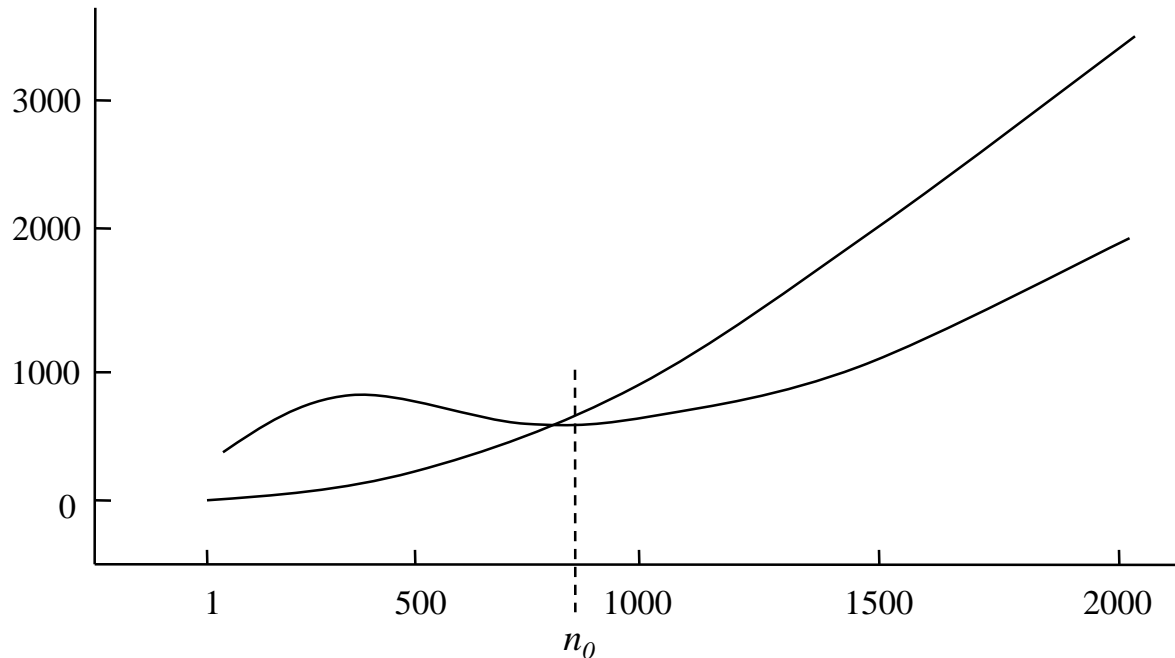
# Best, Worst, and Average Cases

---

- ❖ For some algorithms, different input data of a given size requires different amounts of time.
  - ◆ e.g., to sequential search for an item  $s$  in an array
    - best case:  $s$  is the first element in the array
    - worst case:  $s$  is the last element in the array
    - average case: need to go halfway on average to find  $s$
- ❖ While average time seems to be the fairest measure, it may be difficult to determine.
- ❖ In real-time situations, the worst case time is important.

# Formal Definition of O-Notation

$f(n)$  is  $O(g(n))$  if and only if  
there exists a positive constant  $K$  and  $n_0$ ,  
such that  $|f(n)| \leq K |g(n)|$  for all  $n \geq n_0$ .



The  $f(n)$  curve must  
eventually fit under  
the  $K \cdot g(n)$  curve.

Figure 2: Graphical meaning of O-notation

# Formal Proof of O-Notation

- ❖ The two simplification rules can be proven using the definition.

e.g.,  $f(n) = \frac{3n(n+1)}{2}$  is  $O(n^2)$ .

**Proof:** Let  $K=3$ ,  $n_0=1$ ,  $g(n) = n^2$ .

Substitute in the formal O-notation definition, we have

$|f(n)| \leq K |g(n)|$  for all  $n \geq n_0$ , because

$$\frac{3n(n+1)}{2} \leq 3n^2 \quad (n \geq 1)$$

# Practical Shortcuts for Manipulating O-Notation

❖ An easy way to determine the O-notation for  $f(n)$ .

◆ Separate  $f(n)$  into a dominant term and lesser terms:

$$f(n) = (\text{dominant term}) \pm (\text{lesser terms})$$

◆ Throw away the lesser terms:

$$\begin{aligned} O(f(n)) &= O(\text{dominant term} \pm \text{lesser terms}) \\ &= O(\text{dominant term}) \end{aligned}$$

◆ Ignore and drop coefficients involved.

$$\text{e.g., } O(6n^3 - 15n^2 + 3n \log n) = O(6n^3) = O(n^3)$$

$$O(1) < O(\log n) < O(n) < O(n \log n) < O(n^2) < O(n^3) < O(2^n) < O(10^n)$$

# Remarks on O-Notation

---

- ❖ When writing O-notation, never use the bases of logarithms.
  - ◆ e.g.,  $O(\log n)$ ,  $O(n \log n)$ , but not  $O(\log_2 n)$ ,  $O(n \log_2 n)$ .
  - ◆ **Reason:** changing bases only involves multiplying by new constants (e.g.,  $\log_{10} n = 0.3010 * \log_2 n$ ); and O-notation ignores constants of proportionality.
- ❖ Always use  $O(1)$  to denote the constant O-notation.
  - ◆ **Reason:** for an algorithm running in a number of steps  $f(n)$  which is always less than  $K$ , we have  $f(n) \leq K * 1$ , where  $g(n) = 1$ .

# Remarks on O-Notation (*cont.*)

---

- ❖ When determining the O-notation for  $f(n)$ , try to
  - ◆ make the bound as tight as possible;
    - e.g., if “ $f(n)$  is  $O(n)$ ”, it is also  $O(n^2)$ ,  $O(n^3)$ , ..., but we use only the tightest upper bound.
  - ◆ make the function as simple as possible.
    - Let  $g(n)$  be a single term with a coefficient of one, e.g., we say “ $f(n)$  is  $O(n^2)$ ”, rather than “ $f(n)$  is  $O(3n^2 + n \log n)$ ”.

# When O-Notation Can't Be Trusted?

---

- ❖ Comparisons based on O-notation only apply to large problem sizes.
  - ◆ For small problem sizes, the constants in the running time equations dominate the observed running times.
  - ◆ e.g., to find an item in an ordered list,
    - *sequential search*, *binary search*, and *interpolation search* are  $O(n)$ ,  $O(\log n)$ , and  $O(\log \log n)$ , respectively.
    - when tables contain items -- 1) over 500, *interpolation search* is the fastest; 2) between 20 and 500, *binary search* is the fastest; 3) up to 20, *sequential search* is the fastest.
- ❖ The “measurement and tuning” method is useful to identify optimal solutions to small-sized problems.



# Summary

---

- ❖ It is important to analyze the performance of algorithms and data structures.
- ❖ The problem of measurement:
  - ◆ the hard performance measures (e.g., execution time) vary significantly with different computers, programming languages, and compilers.
- ❖ The O-notation solution:
  - ◆ compare performance based on resource consumption patterns.
  - ◆ intuition, definition, manipulation, and limitations.

# Types of Query Optimization

## ❖ Exhaustive search

- ◆ Cost-based
- ◆ Optimal
- ◆ Combinatorial complexity in the number of relations
- ◆ Workable for **small** solution spaces

## ❖ Heuristics

- ◆ Not optimal
- ◆ Re-group common sub-expressions
- ◆ Perform selection and projection (  $\sigma, \pi$  ) first
- ◆ Replace a join by a series of semijoins
- ◆ Reorder operations to reduce intermediate relation size
- ◆ Optimize individual operations

R  S

# Query Optimization Granularity

---

- ❖ Single query at a time

- ◆ Cannot use common intermediate results

- ❖ Multiple queries at a time

- ◆ Efficient if many similar queries
- ◆ Decision space is much larger

# Query Optimization Timing

---

## ❖ Static

- ◆ Do it **at compilation time** by using statistics, appropriate for exhaustive search, optimized once, but executed many times.
- ◆ Difficult to estimate the size of the intermediate results
- ◆ Can amortize over many executions

## ❖ Dynamic

- ◆ Do it **at execution time**, accurate about the size of the intermediate results, repeated for every execution, expensive.

# Query Optimization Timing (*cont.*)

---

## ❖ Hybrid

- ◆ Compile using a static algorithm
- ◆ If the error in estimate size  $>$  threshold, re-optimize at run time

# Statistics

---

## ❖ Relation

- ◆ Cardinality
- ◆ Size of a tuple
- ◆ Fraction of tuples participating in a join with another relation

## ❖ Attributes

- ◆ Cardinality of the domain
- ◆ Actual number of distinct values

## ❖ Common assumptions

- ◆ Independence between different attribute values
- ◆ Uniform distribution of attribute values within their domain

# Decision Sites

---

- ❖ For query optimization, it may be done by
  - ◆ **Single site** – *centralized* approach
    - Single site determines the best schedule
    - Simple
    - Need knowledge about the entire distributed database
  - ◆ **All the sites involved** – *distributed* approach
    - Cooperation among sites to determine the schedule
    - Need only local information
    - Cost of operation
  - ◆ **Hybrid** – one site makes major decision in cooperation with other sites making local decisions
    - One site determines the global schedule
    - Each site optimizes the local subqueries

# Network Topology

---

- ❖ Wide Area Network (WAN) – point-to-point
  - ◆ Characteristics
    - Low bandwidth
    - Low speed
    - High protocol overhead
  - ◆ Communication cost will dominate; ignore all other cost factors
  - ◆ Global schedule to minimize communication cost
  - ◆ Local schedules according to centralized query optimization



# Network Topology (*cont.*)

---

## ❖ Local Area Network (LAN)

- ◆ Communication cost not that dominate
- ◆ Total cost function should be considered
- ◆ Broadcasting can be exploited
- ◆ Special algorithms exist for star networks

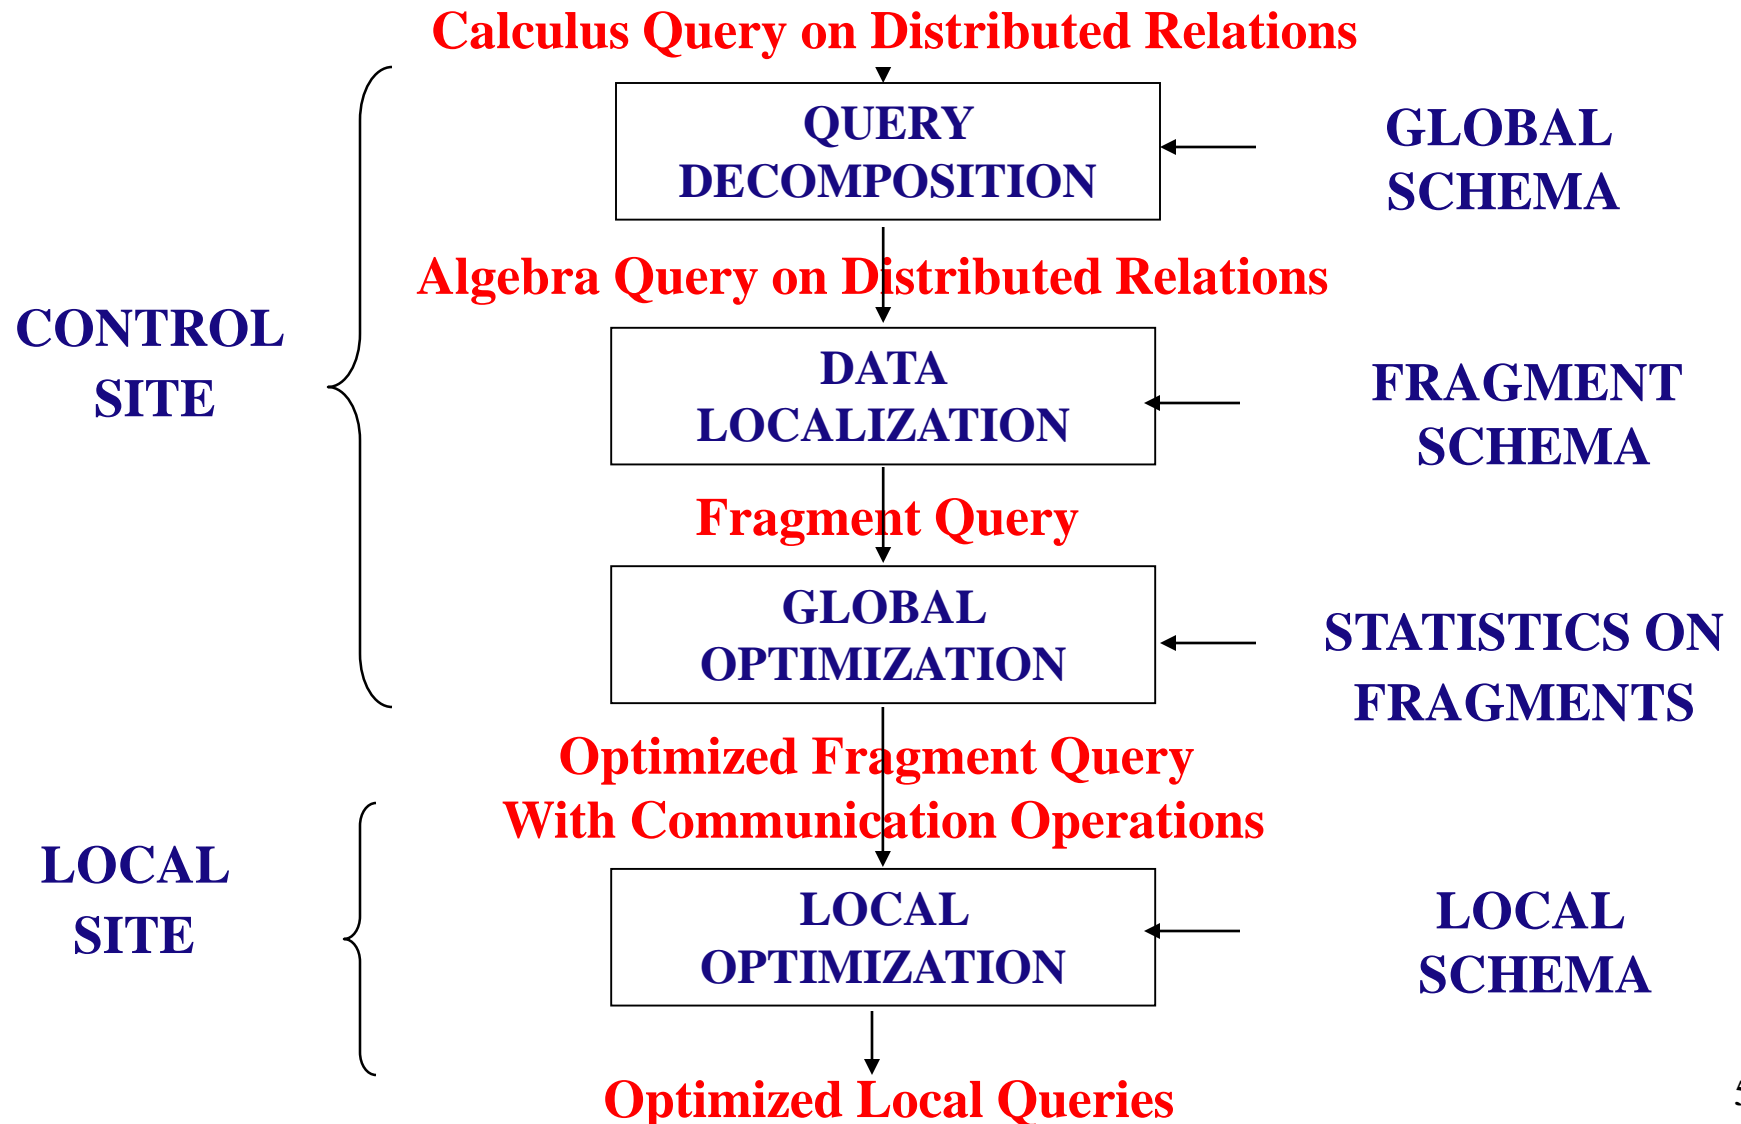
# Other Information to Exploit

---

- ❖ Using **replications** to minimize communication costs
- ❖ Using **semijoins** to reduce the size of operand relations to cut down communication costs when overhead is not significant.

R  S

# Layers of Query Processing



# Step 1 - Query Decomposition

---

- ❖ Decompose calculus query into algebra query using global conceptual schema information.

(1) normalization



(2) analysis



(3) elimination of redundancy



(4) rewriting

# Step 1 - Query Decomposition (*cont.*)

---

## 1) Normalization

- ♦ The calculus query is written in a normalized form (CNF or DNF) for subsequent manipulation

## 2) Analysis

- ♦ To reject normalized queries for which further processing is either impossible or unnecessary (type incorrect or semantically incorrect)

## 3) Simplification (elimination of redundancy)

- ♦ Redundant predicates are eliminated to obtain simplified queries

## 4) Rewriting

- ♦ The calculus query is translated to optimal algebraic query representation
- ♦ More than one translation is possible

# 1) Normalization

## ❖ Lexical and syntactic analysis

- ◆ check validity (similar to compilers)
- ◆ check for attributes and relations
- ◆ type checking on the qualification

## ❖ There are two possible forms of representing the predicates in query qualification

- ◆ Conjunctive Normal Form (CNF) or Disjunctive Normal Form (DNF)
  - CNF:  $(p_{11} \vee p_{12} \vee \dots \vee p_{1n}) \wedge \dots \wedge (p_{m1} \vee p_{m2} \vee \dots \vee p_{mn})$
  - DNF:  $(p_{11} \wedge p_{12} \wedge \dots \wedge p_{1n}) \vee \dots \vee (p_{m1} \wedge p_{m2} \wedge \dots \wedge p_{mn})$
  - OR's mapped into union
  - AND's mapped into join or selection

# 1) Normalization (*cont.*)

- ❖ The transformation of the quantifier-free predicate is straightforward using the well-known equivalence rules for logical operations ( $\wedge \vee \neg$ )

$$P_1 \wedge P_2 \Leftrightarrow P_2 \wedge P_1$$

$$P_1 \vee P_2 \Leftrightarrow P_2 \vee P_1$$

$$P_1 \wedge (P_2 \wedge P_3) \Leftrightarrow (P_1 \wedge P_2) \wedge P_3$$

$$P_1 \vee (P_2 \vee P_3) \Leftrightarrow (P_1 \vee P_2) \vee P_3$$

$$P_1 \wedge (P_2 \wedge P_3) \Leftrightarrow (P_1 \wedge P_2) \wedge P_3$$

$$P_1 \vee (P_2 \wedge P_3) \Leftrightarrow (P_1 \vee P_2) \wedge (P_1 \vee P_3)$$

$$P_1 \wedge (P_2 \vee P_3) \Leftrightarrow (P_1 \wedge P_2) \vee (P_1 \wedge P_3)$$

$$\neg(P_1 \wedge P_2) \Leftrightarrow \neg P_1 \vee \neg P_2$$

$$\neg(P_1 \vee P_2) \Leftrightarrow \neg P_1 \wedge \neg P_2$$

$$\neg(\neg P_1) \Leftrightarrow P_1$$

# 1) Normalization (*cont.*)

## ❖ Example

```
SELECT  ENAME
FROM    EMP, ASG
WHERE    EMP.ENO=ASG.ENO AND ASG.JNO="J1"
           AND (DUR=12 OR DUR=24)
```

## ❖ The conjunctive normal form:

$$\begin{aligned} &EMP.ENO = ASG.ENO \\ &\wedge ASG.JNO = "J1" \\ &\wedge (DUR = 12 \vee DUR = 24) \end{aligned}$$



## 2) Analysis

---

### ❖ Objective

- ◆ reject **type incorrect** or **semantically incorrect** queries

### ❖ Type incorrect

- ◆ if any of its attribute or relation names is not defined in the global schema
- ◆ if operations are applied to attributes of the wrong type

## 2) Analysis (*cont.*)

❖ Type incorrect example

**SELECT**     E #

**FROM**       EMP

**WHERE**      ENAME > 200

! Undefined  
attribute

! Type  
mismatch

## 2) Analysis (*cont.*)

---

### ❖ Semantically incorrect

- ◆ Components do not contribute in any way to the generation of the result
- ◆ For only those queries that do not use disjunction ( $\vee$ ) or negation ( $\neg$ ), semantic correctness can be determined by using *query graph*

# Query Graph

---

## ❖ Two kinds of nodes

- ◆ One node represents the result relation
- ◆ Other nodes represent operand relations

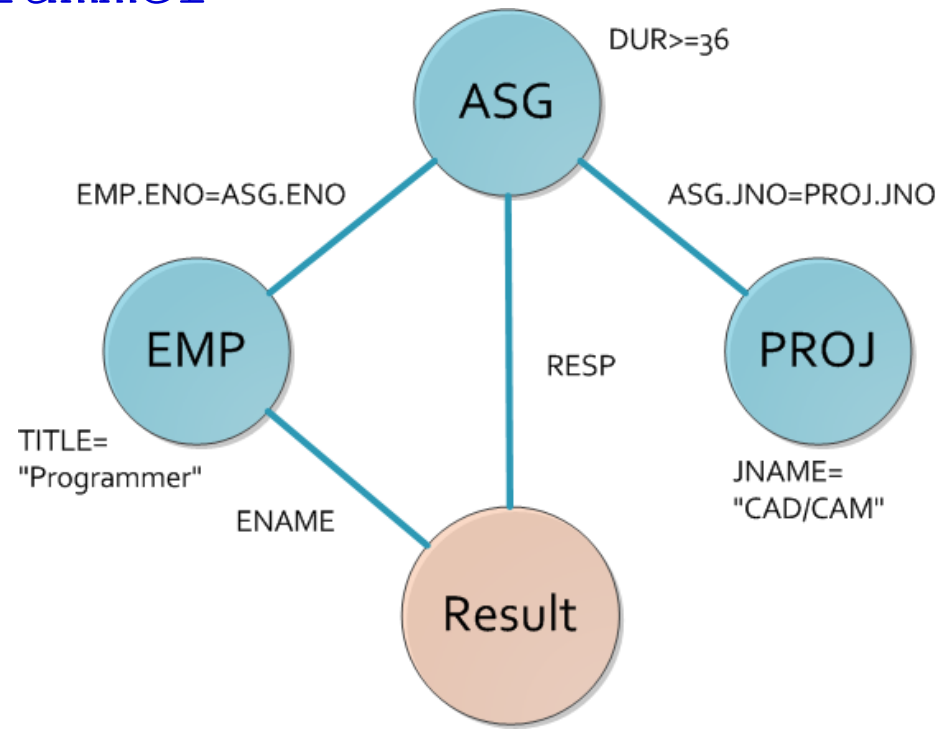
## ❖ Two types of edges

- ◆ an edge to represent a join if neither of its two nodes is the result
- ◆ an edge to represent a projection if one of its node is the result node

Nodes and edges may be labeled by predicates for selection, projection, or join.

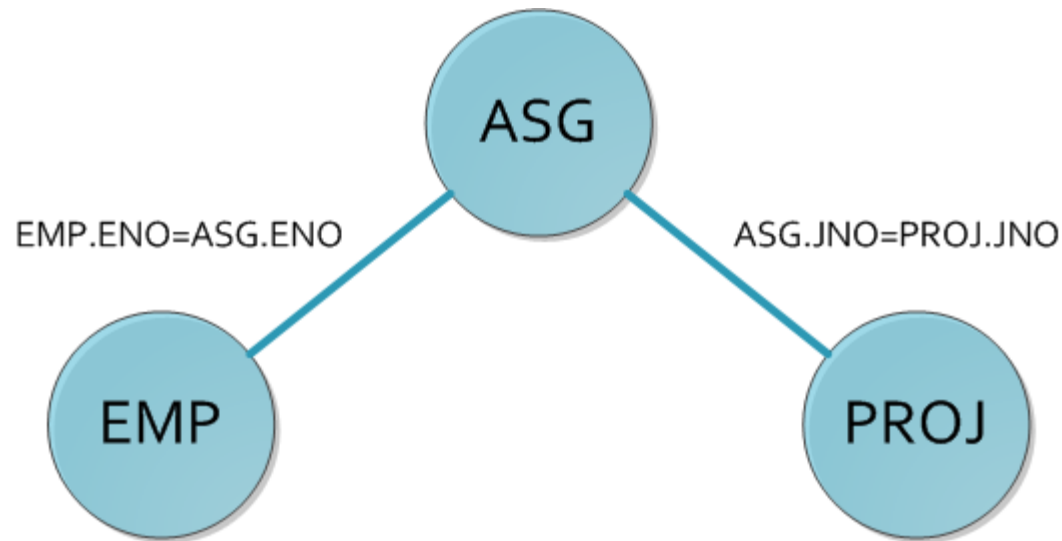
# Query Graph Example

```
SELECT  ENAME, RESP
FROM    EMP, ASG, PROJ
WHERE    EMP.ENO=ASG.GNO AND ASG.PNO=PROJ.PNO
AND      PNAME="CAD/CAM" AND DUR>36
AND      TITLE="Programmer"
```



# Join Graph Example 1

A **subgraph** of query graph for join operation.



# Tool of Analysis

---

- ❖ A conjunctive query **without negation** is semantically incorrect if its query graph is **NOT** connected!

# Analysis Example

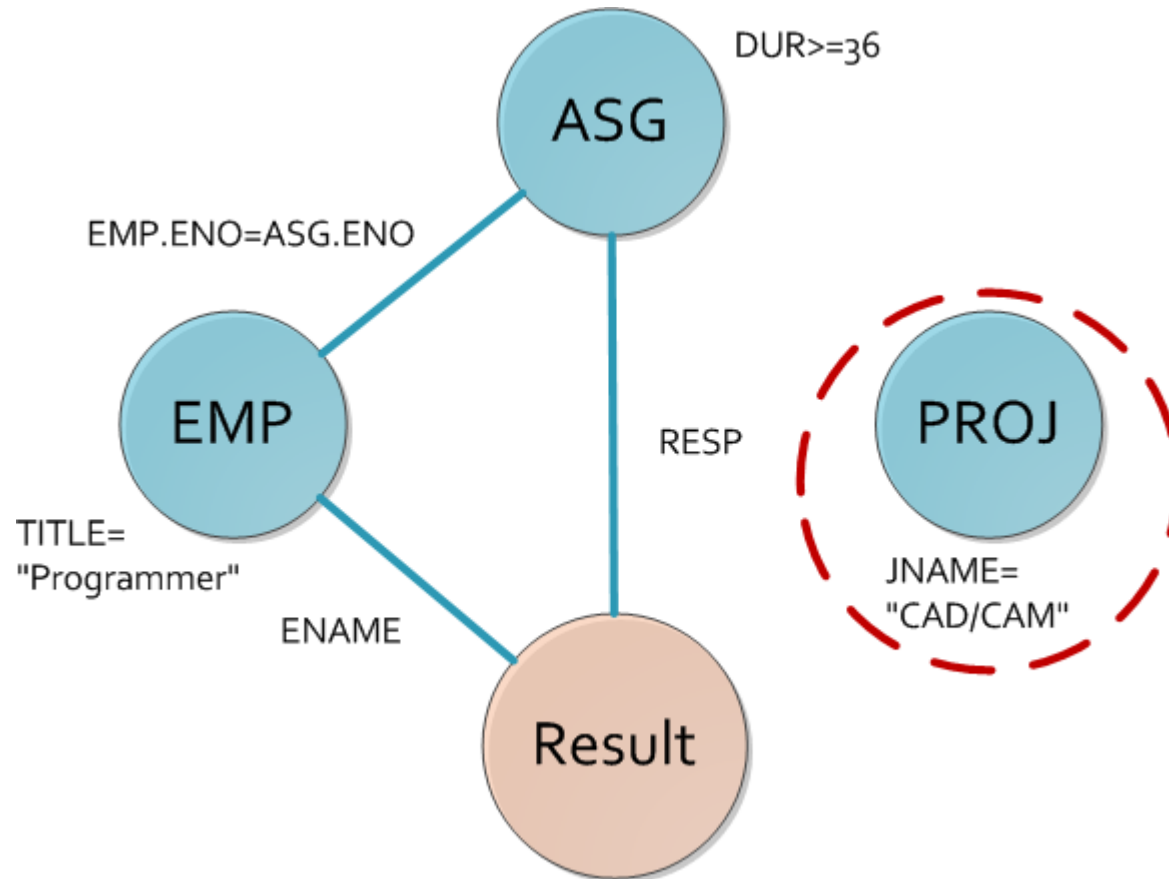
---

## Example 2

```
SELECT      ENAME, RESP
FROM        EMP, ASG, PROJ
WHERE       EMP.ENO=ASG.GNO
AND       ASG.PNO=PROJ.PNO
AND        PNAME="CAD/CAM"
AND        DUR>36
AND        TITLE="Programmer"
```



# Query Graph Example 2



### 3) Simplification

❖ Using idempotency rules to eliminate redundant predicates from WHERE clause.

$$P \wedge P \Leftrightarrow P$$

$$P \vee P \Leftrightarrow P$$

$$P \wedge \text{true} \Leftrightarrow P$$

$$P \vee \text{false} \Leftrightarrow P$$

$$P \wedge \text{false} \Leftrightarrow \text{false}$$

$$P \vee \text{true} \Leftrightarrow \text{true}$$

$$P \wedge \neg P \Leftrightarrow \text{false}$$

$$P \vee \neg P \Leftrightarrow \text{true}$$

$$P_1 \wedge (P_1 \vee P_2) \Leftrightarrow P_1$$

$$P_1 \vee (P_1 \wedge P_2) \Leftrightarrow P_1$$

# Simplification Example

```
SELECT  TITLE
FROM    EMP
WHERE   (NOT (TITLE="Programmer")
        AND (TITLE="Programmer" OR
              TITLE="Electrical Eng."))
        AND NOT (TITLE="Electrical Eng.")
        OR  ENAME="J.Doe"
```

p1 = <TITLE = ``Programmer">  
p2 = <TITLE = ``Elec. Engr">  
p3 = <ENAME = ``J.Doe">

Let the query qualification is  
 $(\neg p1 \wedge (p1 \vee p2) \wedge \neg p2) \vee p3$

The disjunctive normal form of the query is  
$$= (\neg p1 \wedge p1 \wedge \neg p2) \vee (\neg p1 \wedge p2 \wedge \neg p2) \vee p3$$
$$= (false \wedge \neg p2) \vee (\neg p1 \wedge false) \vee p3$$
$$= false \vee false \vee p3$$
$$= p3$$

# Simplification Example

---

```
SELECT  TITLE
FROM    EMP
WHERE    (NOT (TITLE="Programmer")
           AND  (TITLE="Programmer"
           OR    TITLE="Electrical Eng.")
           AND NOT (TITLE="Electrical Eng."))
           OR  ENAME="J.Doe"
```

is equivalent to

```
SELECT  TITLE
FROM    EMP
WHERE    ENAME="J.Doe"
```

# 4) Rewriting

---

- ❖ Converting a calculus query in relational algebra
  - ♦ straightforward transformation from relational calculus to relational algebra
  - ♦ restructuring relational algebra expression to improve performance
  - ♦ making use of query trees

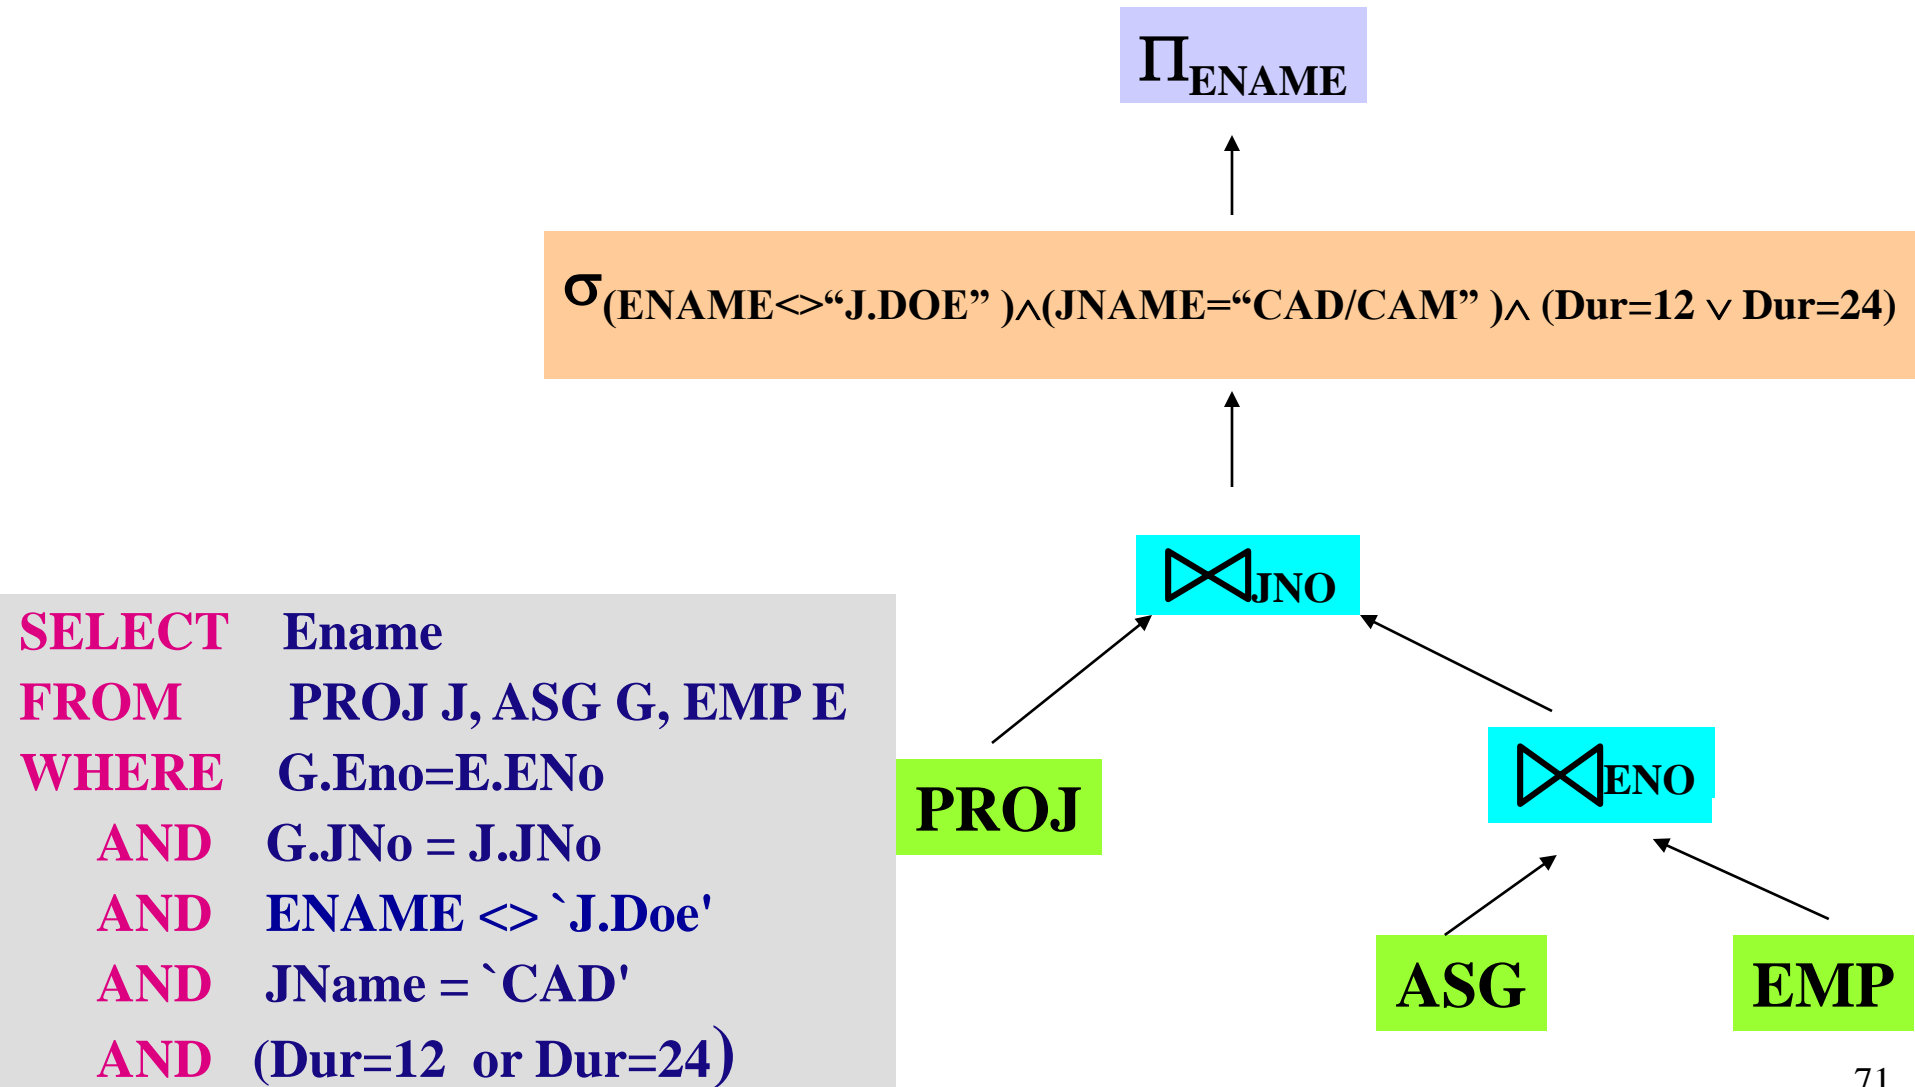
# Relational Algebra Tree

---

## ❖ A tree defined by:

- ◆ a root node representing the query result
- ◆ leaves representing database relations
- ◆ non-leaf nodes representing relations produced by operations
- ◆ edges from leaves to root representing the sequences of operations

# An SQL Query and Its Query Tree



# How to translate an SQL query into an algebra tree?

---

1. Create a leaf for every relation in the FROM clause
2. Create the root as a project operation involving attributes in the SELECT clause
3. Create the operation sequence by the predicates and operators in the WHERE clause



# Rewriting -- Transformation Rules (I)

- ❖ Commutativity of binary operations:

$$\mathbf{R} \times \mathbf{S} \Leftrightarrow \mathbf{S} \times \mathbf{R}$$

$$\mathbf{R} \bowtie \mathbf{S} \Leftrightarrow \mathbf{S} \bowtie \mathbf{R}$$

$$\mathbf{R} \cup \mathbf{S} \Leftrightarrow \mathbf{S} \cup \mathbf{R}$$

- ❖ Associativity of binary operations:

$$(\mathbf{R} \times \mathbf{S}) \times \mathbf{T} \Leftrightarrow \mathbf{R} \times (\mathbf{S} \times \mathbf{T})$$

$$(\mathbf{R} \bowtie \mathbf{S}) \bowtie \mathbf{T} \Leftrightarrow \mathbf{R} \bowtie (\mathbf{S} \bowtie \mathbf{T})$$

- ❖ Idempotence of unary operations: grouping of projections and selections

- ♦  $\Pi_{A'} (\Pi_{A''} (\mathbf{R})) \Leftrightarrow \Pi_{A'} (\mathbf{R})$  for  $A' \subseteq A'' \subseteq A$

- ♦  $\sigma_{p1(A1)} (\sigma_{p2(A2)} (\mathbf{R})) \Leftrightarrow \sigma_{p1(A1) \wedge p2(A2)} (\mathbf{R})$

# Rewriting -- Transformation Rules (II)

## ❖ Commuting selection with projection

$$\Pi_{A_1, \dots, A_n} (\sigma_{p(A_p)}(R)) \Leftrightarrow \Pi_{A_1, \dots, A_n} (\sigma_{p(A_p)}(\Pi_{A_1, \dots, A_n, A_p}(R)))$$

## ❖ Commuting selection with binary operations

$$\sigma_{p(A_i)}(R \times S) \Leftrightarrow (\sigma_{p(A_i)}(R)) \times S$$

$$\sigma_{p(A_i)}(R \bowtie S) \Leftrightarrow (\sigma_{p(A_i)}(R)) \bowtie S$$

$$\sigma_{p(A_i)}(R \cup S) \Leftrightarrow \sigma_{p(A_i)}(R) \cup \sigma_{p(A_i)}(S)$$

## ❖ Commuting projection with binary operations

$$\Pi_C(R \times S) \Leftrightarrow \Pi_A(R) \times \Pi_B(S) \quad \text{where } C = A \cup B$$

$$\Pi_C(R \bowtie S) \Leftrightarrow \Pi_C(R) \bowtie \Pi_C(S)$$

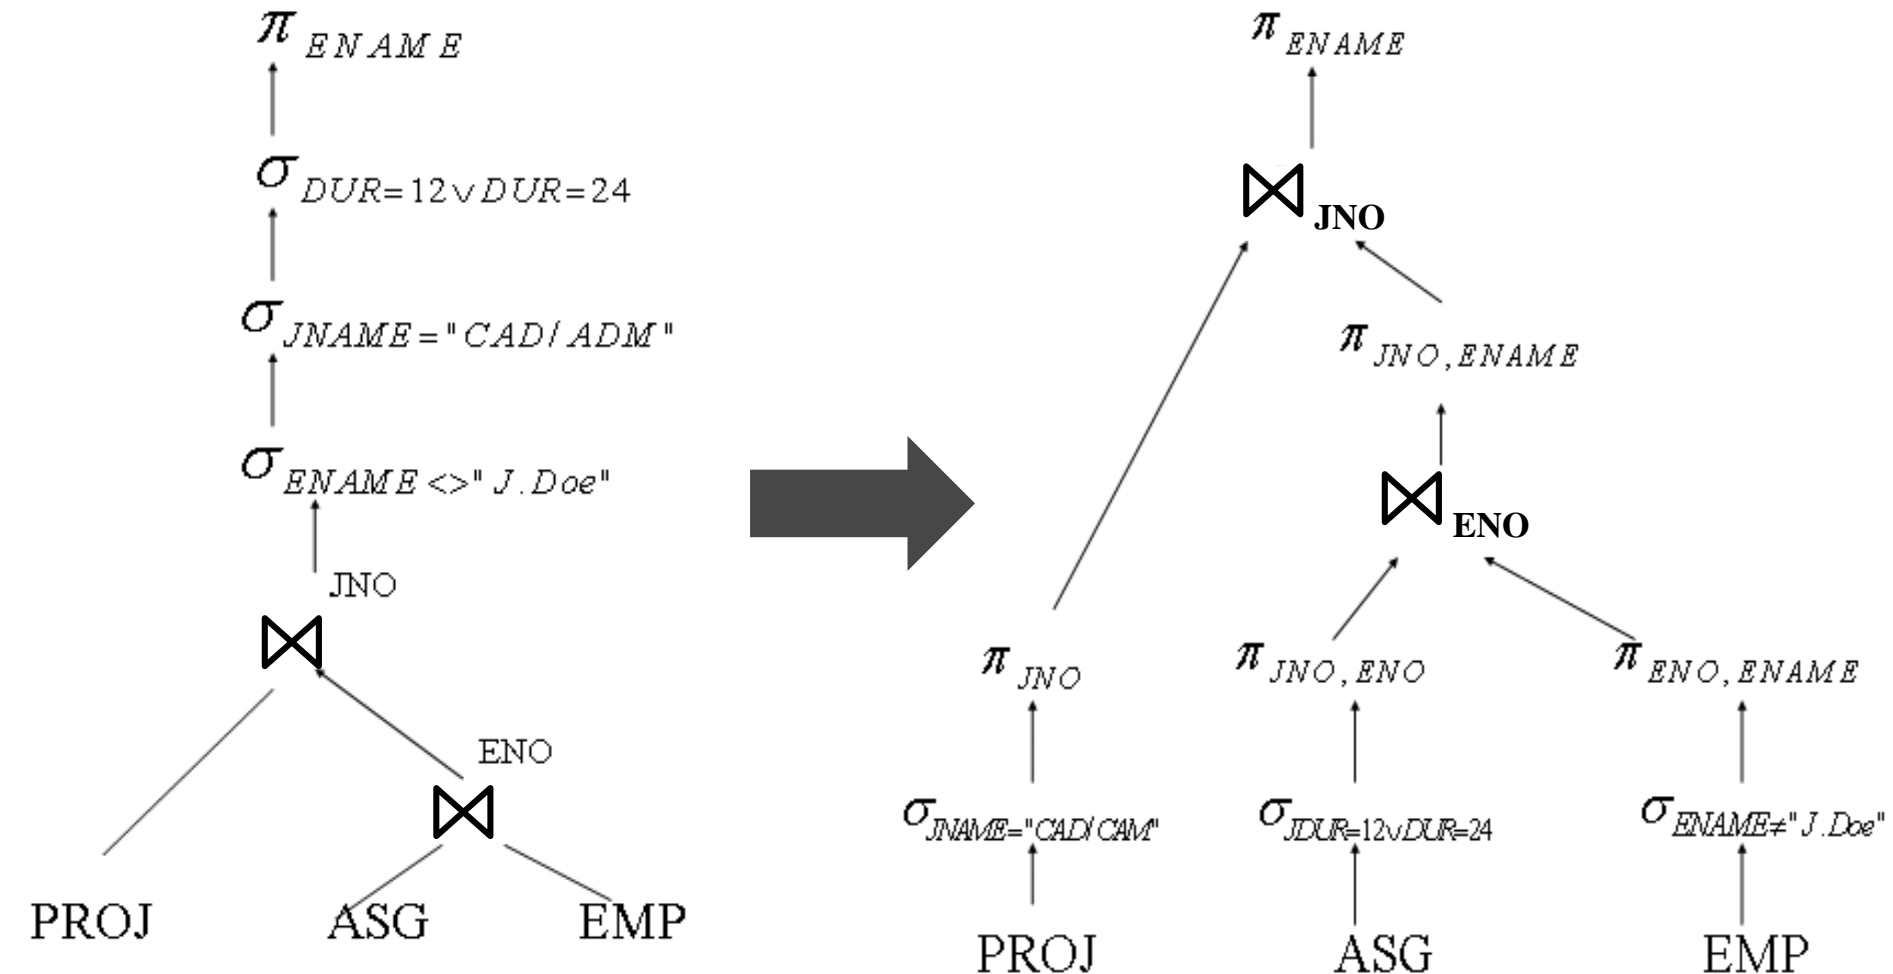
$$\Pi_C(R \cup S) \Leftrightarrow \Pi_C(R) \cup \Pi_C(S)$$

# How to use transformation rules to optimize?

---

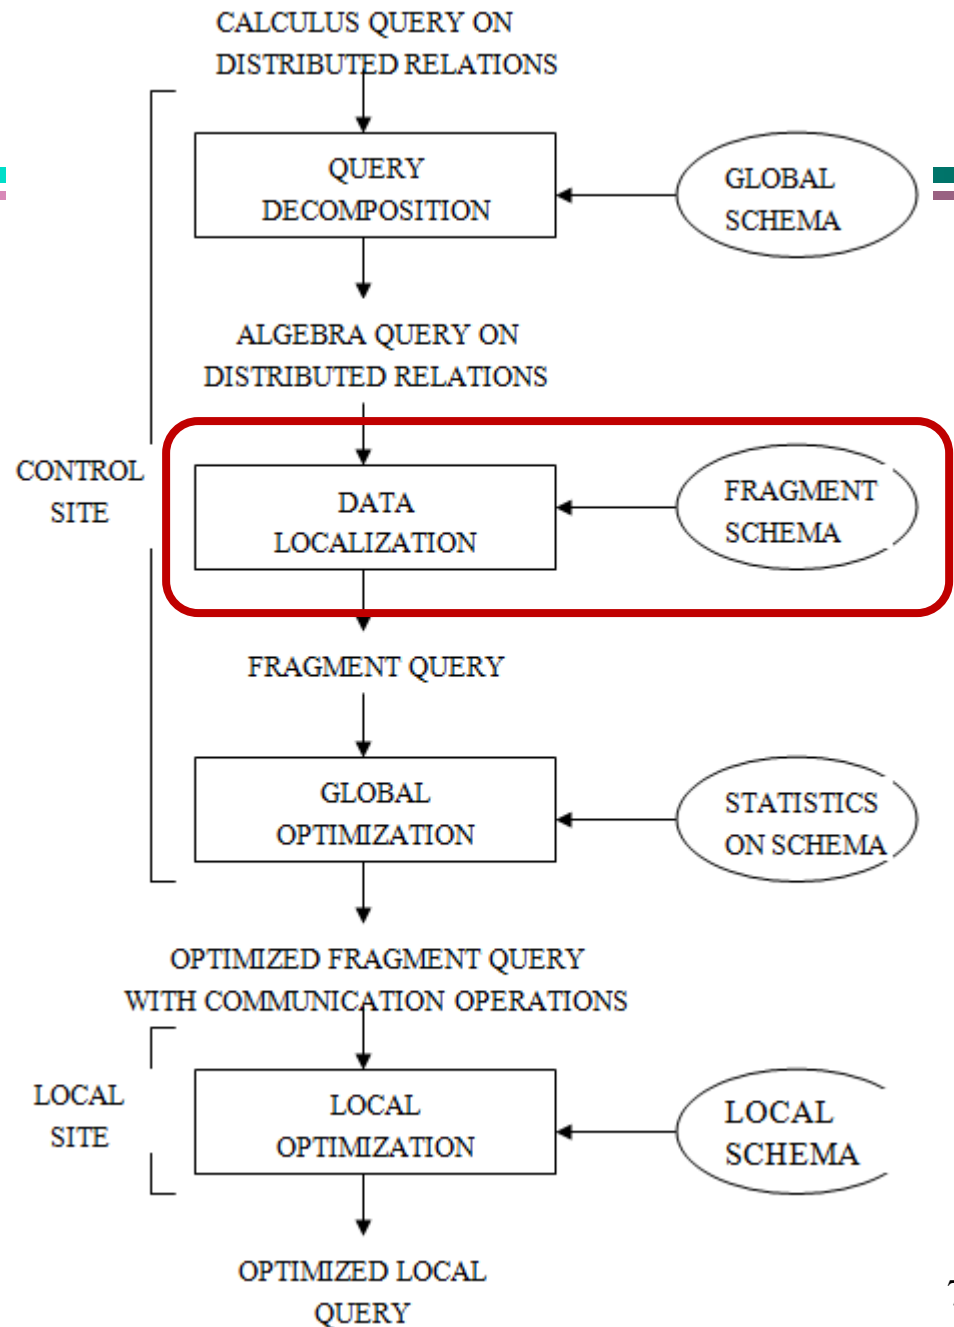
- ❖ Unary operations on the same relation may be grouped to access the same relation once
- ❖ Unary operations may be commuted with binary operations, so that they may be performed first to reduce the size of intermediate relations
- ❖ Binary operations may be reordered

# Optimization of Previous Query Tree



# Step 2 : Data Localization

- ❖ Task : To translate a query on global relation into algebra queries on physical fragments, and optimize the query by reduction.



# Reduction with Selection for PHF

```
SELECT *  
FROM EMP  
WHERE ENO="E5"
```

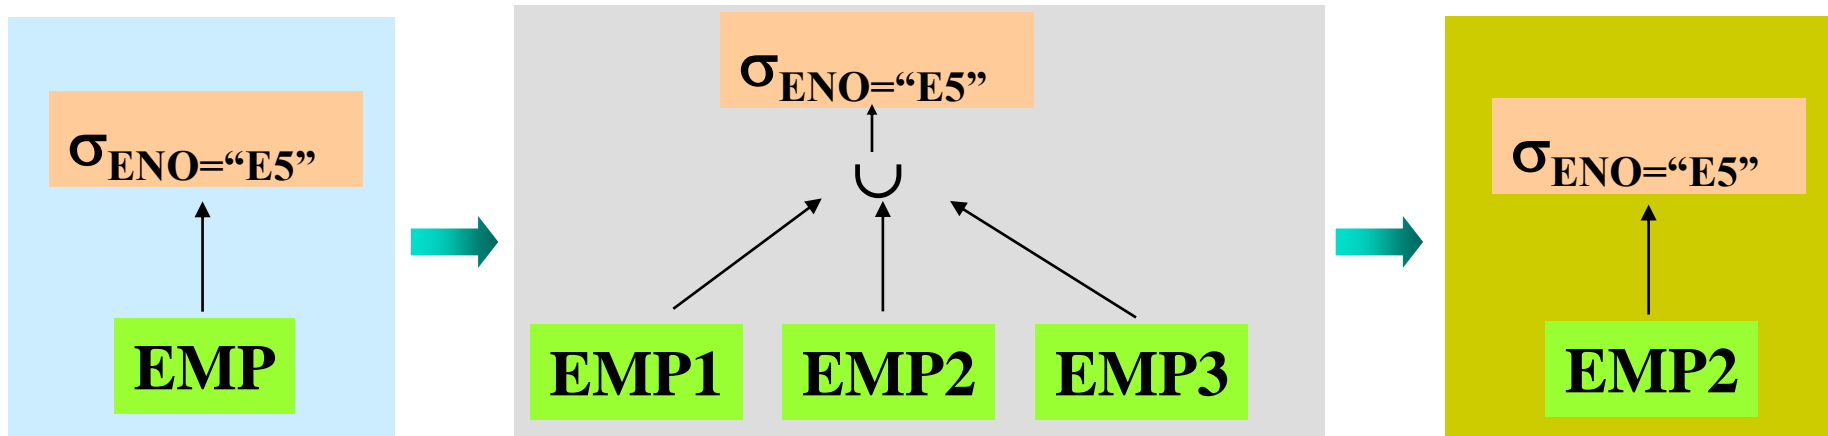
EMP is fragmented into

$EMP1 = \sigma_{ENO \leq "E3"}(EMP)$

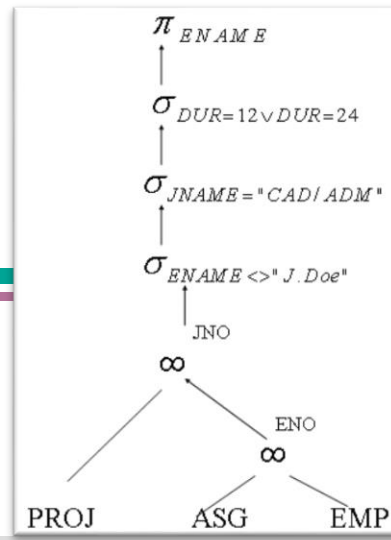
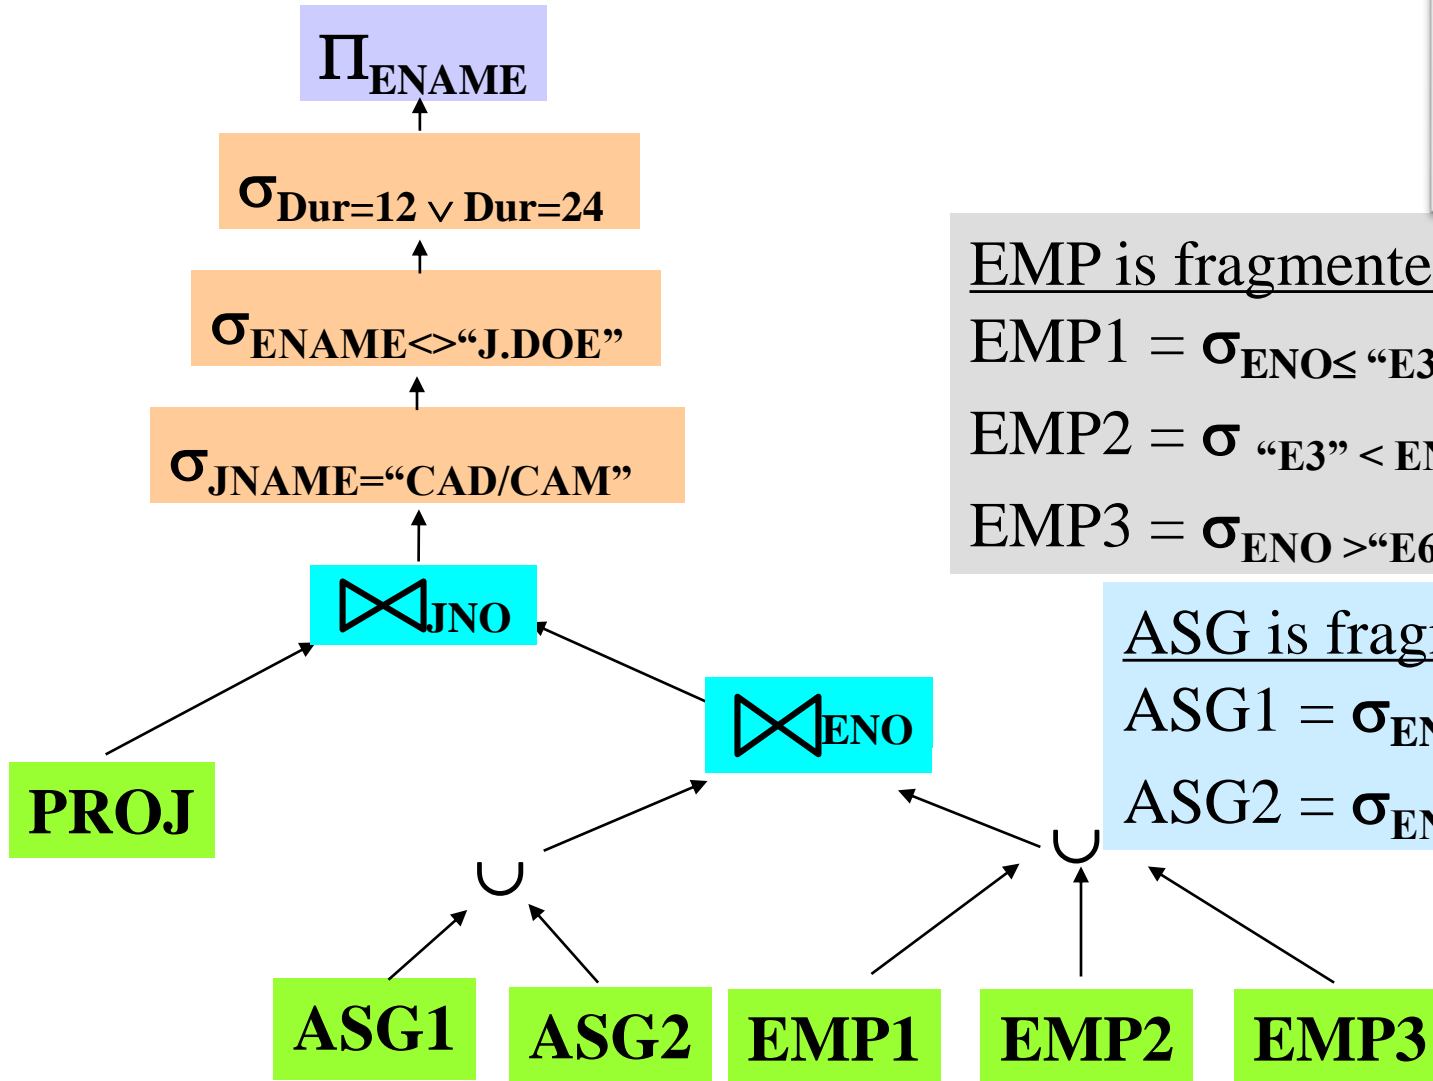
$EMP2 = \sigma_{"E3" < ENO \leq "E6"}(EMP)$

$EMP3 = \sigma_{ENO > "E6"}(EMP)$

Given Relation  $R$ ,  $F_R = \{R_1, R_2, \dots, R_n\}$  where  $R_j = \sigma_{p_j}(R)$   
 $\sigma_{p_j}(R_j) = \emptyset$  if  $\forall x \in R: \neg(p_i(x) \wedge p_j(x))$



# Data Localization - Example



EMP is fragmented into

EMP1 =  $\sigma_{ENO \leq "E3"} (EMP)$

EMP2 =  $\sigma_{"E3" < ENO \leq "E6"} (EMP)$

EMP3 =  $\sigma_{ENO > "E6"} (EMP)$

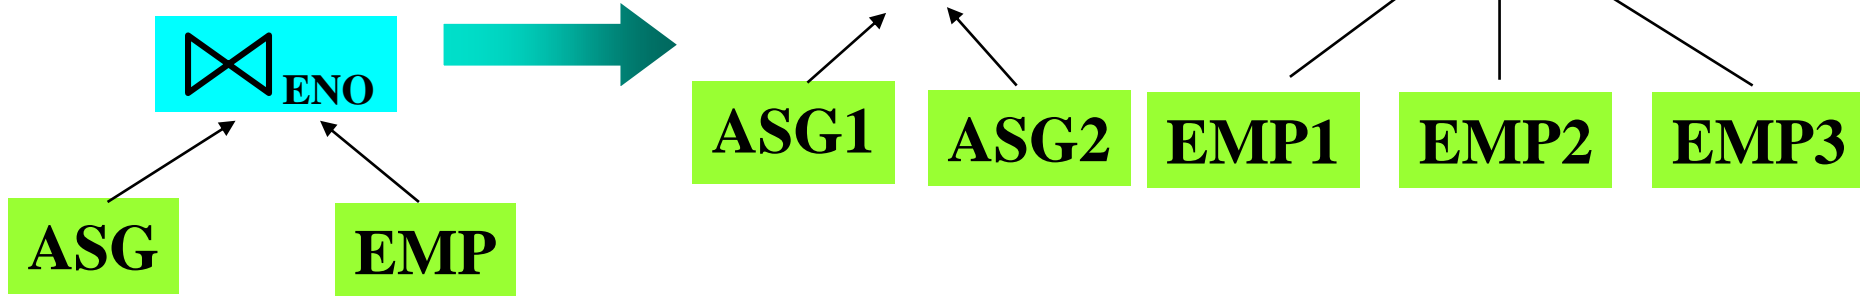
ASG is fragmented into

ASG1 =  $\sigma_{ENO \leq "E3"} (ASG)$

ASG2 =  $\sigma_{ENO > "E3"} (ASG)$

# Reduction with Join for PHF

```
SELECT *  
FROM EMP, ASG  
WHERE EMP.ENO=ASG.ENO
```

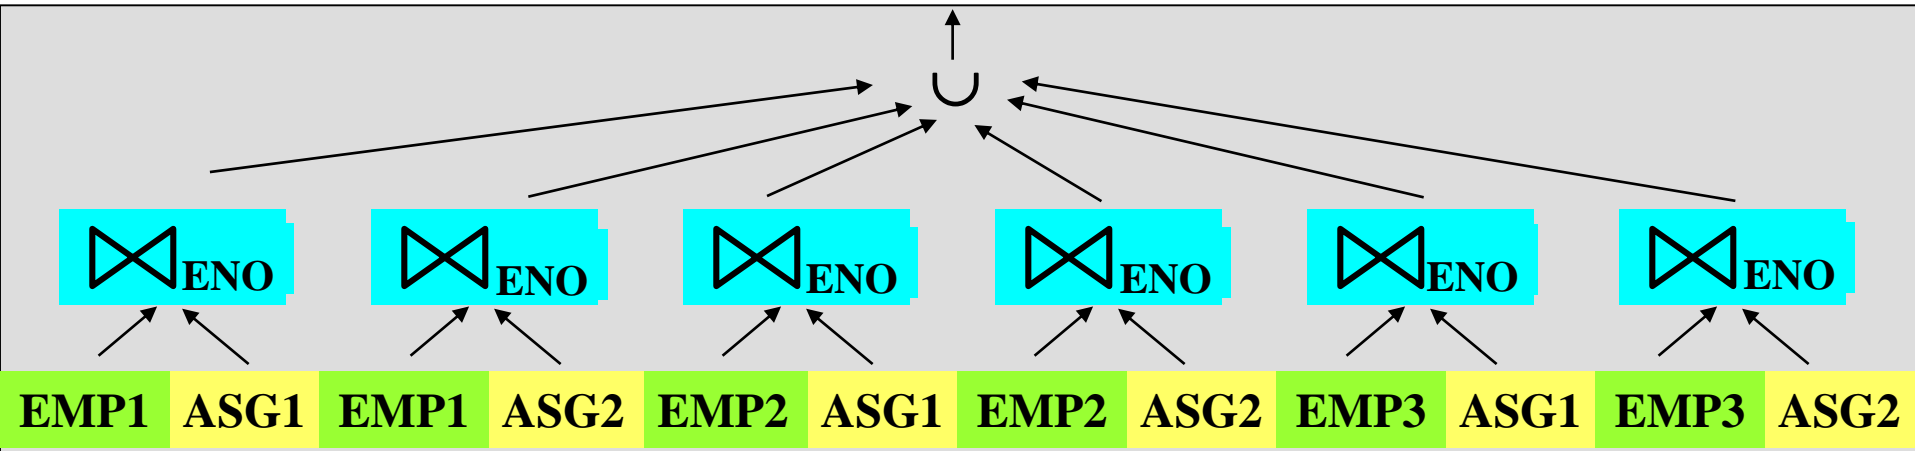
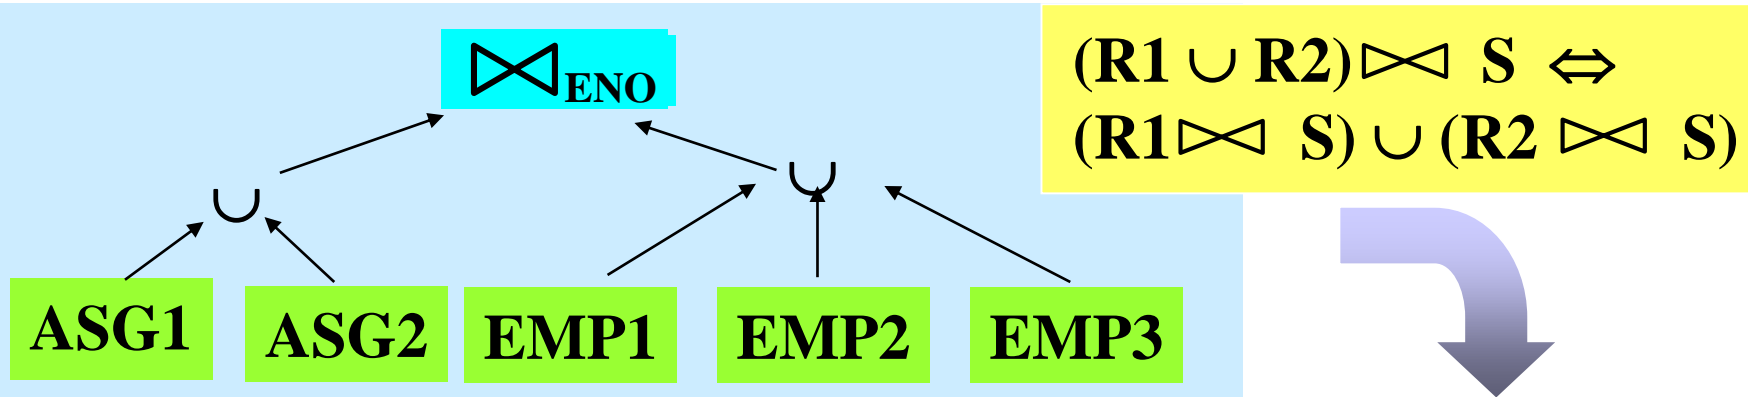


ASG is fragmented into  
 $\text{ASG1} = \sigma_{\text{ENO} \leq \text{"E3"}}(\text{ASG})$   
 $\text{ASG2} = \sigma_{\text{ENO} > \text{"E3"}}(\text{ASG})$

EMP is fragmented into  
 $\text{EMP1} = \sigma_{\text{ENO} \leq \text{"E3"}}(\text{EMP})$   
 $\text{EMP2} = \sigma_{\text{"E3"} < \text{ENO} \leq \text{"E6"}}(\text{EMP})$   
 $\text{EMP3} = \sigma_{\text{ENO} > \text{"E6"}}(\text{EMP})$



# Reduction with Join for PHF (I)



$EMP1 = \sigma_{ENO \leq "E3"}(EMP)$

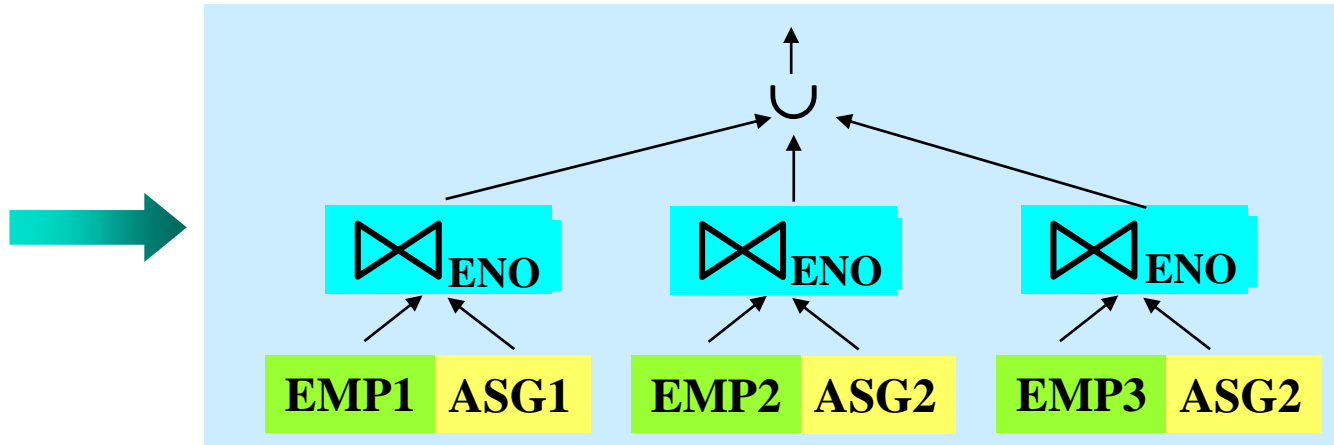
$EMP2 = \sigma_{"E3" < ENO \leq "E6"}(EMP)$

$EMP3 = \sigma_{ENO > "E6"}(EMP)$

$ASG1 = \sigma_{ENO \leq "E3"}(ASG)$

$ASG2 = \sigma_{ENO > "E3"}(ASG)$

# Reduction with Join for PHF (II)



Given  $R_i = \sigma_{p_i}(R)$  and  $R_j = \sigma_{p_j}(R)$

$R_i \bowtie R_j = \emptyset$  if  $\forall x \in R_i, \forall y \in R_j: \neg(p_i(x) \wedge p_j(y))$

**Reduction with join**

1. Distribute join over union
2. Eliminate unnecessary work

# Reduction for VF

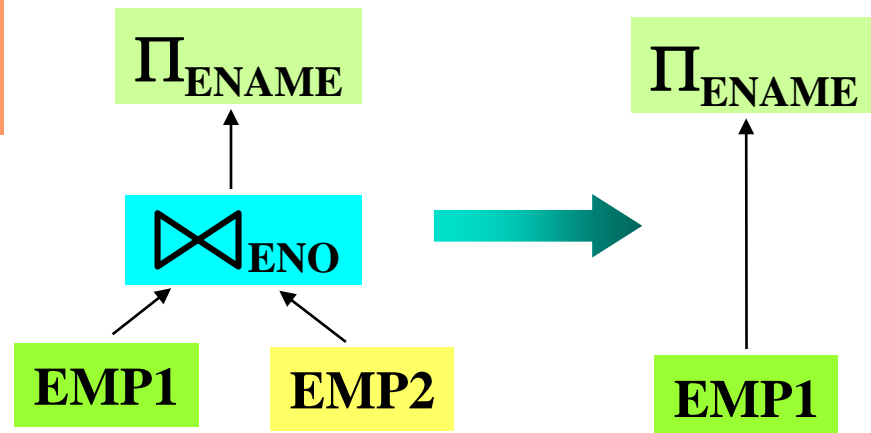
## ❖ Find useless intermediate relations

Relation  $R$  defined over attributes  $A = \{A1, A2, \dots, An\}$  vertically fragmented as  $R_i = \Pi_{A'}(R)$  where  $A' \subseteq A$   
 $\Pi_D(R_i)$  is useless if the set of projection attributes  $D$  is not in  $A'$

**EMP1 =  $\Pi_{ENO,ENAME}(EMP)$**

**EMP2 =  $\Pi_{ENO,TITLE}(EMP)$**

```
SELECT ENAME  
FROM EMP
```



# Reduction for DHF

Distribute joins over union

Apply the join reduction for horizontal fragmentation

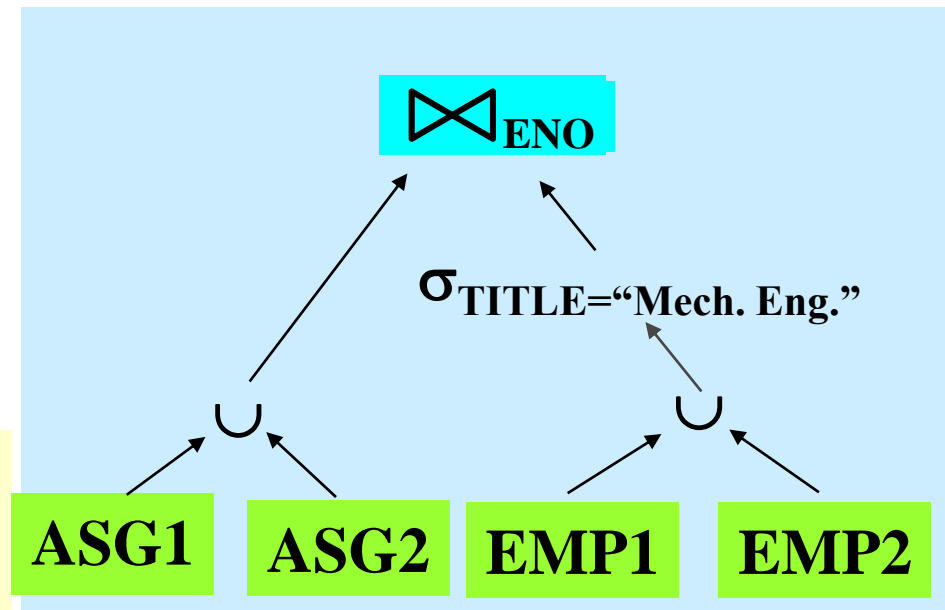
EMP1:  $\sigma_{\text{TITLE}=\text{"Programmer"}}(\text{EMP})$

EMP2:  $\sigma_{\text{TITLE}\neq\text{"Programmer"}}(\text{EMP})$

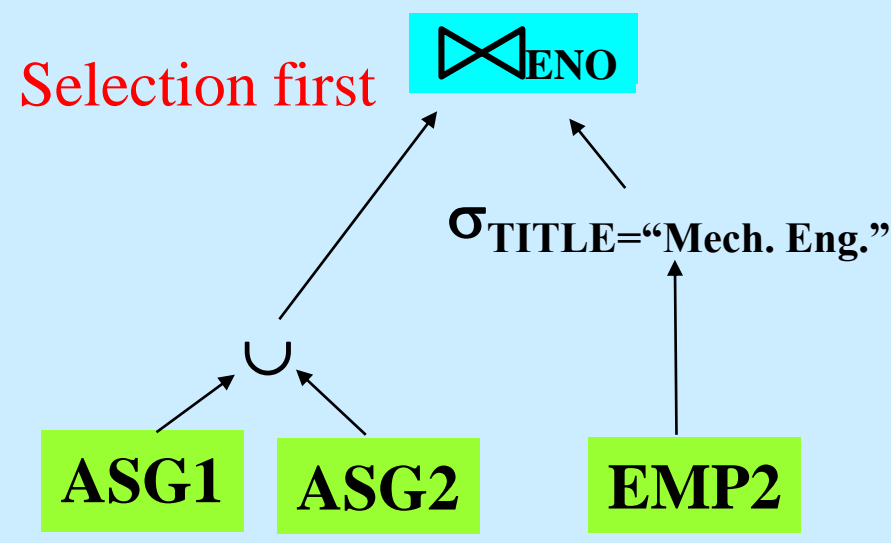
ASG1:  $\text{ASG} \bowtie_{\text{ENO}} \text{EMP1}$

ASG2:  $\text{ASG} \bowtie_{\text{ENO}} \text{EMP2}$

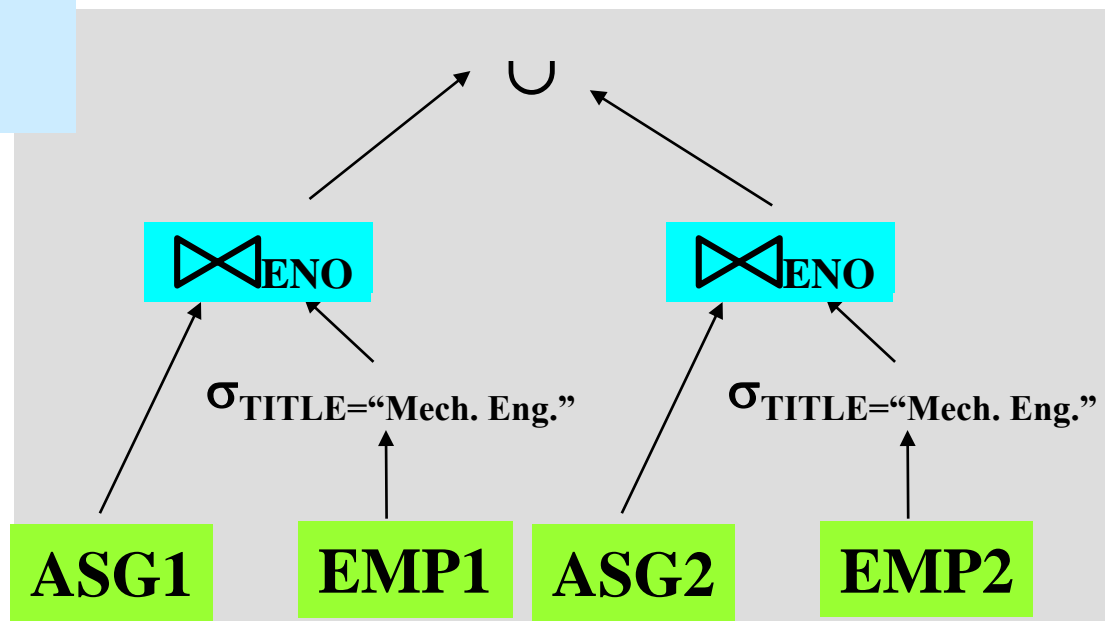
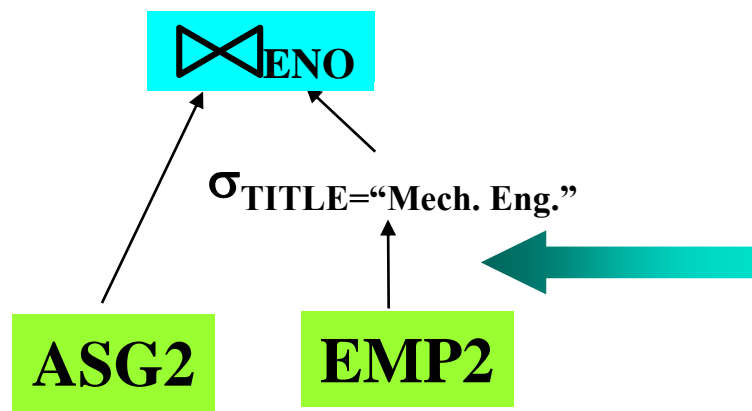
```
SELECT *  
FROM EMP, ASG  
WHERE ASG.ENO = EMP.ENO  
AND EMP.TITLE = "Mech. Eng."
```



# Reduction for DHF (II)



Joins over union



# Reduction for Hybrid Fragmentation

---

- ❖ Combine the rules already specified
  - ◆ Remove *empty relations* generated by contradicting selection on horizontal fragments;
  - ◆ Remove *useless relations* generated by projections on vertical fragments;
  - ◆ Distribute *joins over unions* in order to isolate and remove useless joins.

# Reduction for Hybrid Fragmentation

## - Example

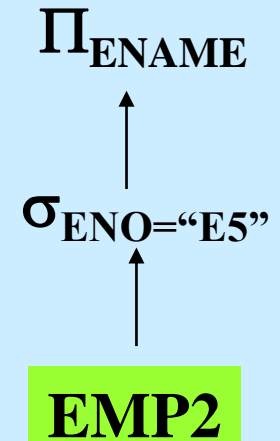
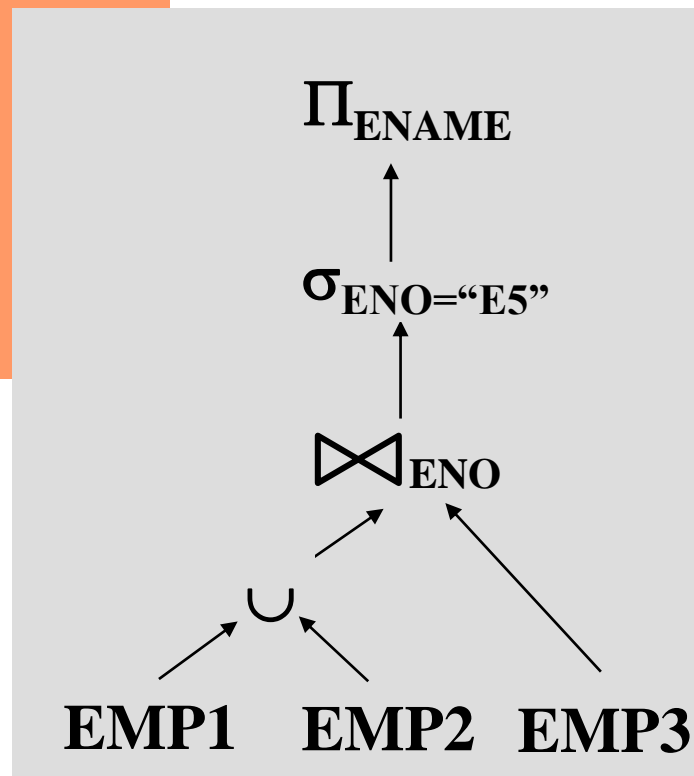
$EMP1 = \sigma_{ENO \leq "E4"} (\Pi_{ENO, ENAME} (EMP))$

$EMP2 = \sigma_{ENO > "E4"} (\Pi_{ENO, ENAME} (EMP))$

$EMP3 = \Pi_{ENO, TITLE} (EMP)$

QUERY:

**SELECT** ENAME  
**FROM** EMP  
**WHERE** ENO = "E5"



---

# Question & Answer